Data Analysis with Python: From Fundamentals to Advanced Techniques

Part I: Foundations of Data Analysis

1. Introduction to Data Analysis for Python Experts

- What is Data Analysis?
- The Role of Python in Modern Data Analysis
- Understanding the Data Analysis Workflow
- o Setting Up a High-Performance Data Analysis Environment

2. Core Concepts in Data Analysis

- o Understanding Structured, Semi-Structured, and Unstructured Data
- o The Data Analysis Pipeline: From Raw Data to Insights
- Exploratory Data Analysis (EDA) vs. Explanatory Analysis
- Data Ethics, Integrity, and Reproducibility

3. Python Essentials for Data Analysis (Quick Recap)

- Working with Iterators, Generators, and List Comprehensions
- Lambda Functions, Map-Reduce, and Functional Programming Techniques
- Error Handling and Debugging for Data Analysis
- Efficient File I/O: Working with CSV, JSON, and SQL

Part II: Mastering Numpy for Numerical Computing

4. Numpy Essentials

- Numpy Arrays: The Foundation of Numerical Computing
- Vectorization: Replacing Loops with Efficient Operations
- Broadcasting, Indexing, and Slicing Techniques
- Statistical Functions, Aggregations, and Basic Linear Algebra

5. Advanced Numpy Techniques

- Memory Management and Data Alignment
- Structured Arrays and Record Arrays
- Optimizing Computations with Numba and Cython
- Random Sampling, Simulations, and Complex Mathematical Operations

Part III: Data Manipulation with Pandas

6. Pandas Fundamentals

- Series, DataFrames, and MultiIndex Structures
- Data Import/Export (CSV, Excel, JSON, SQL, APIs)
- Indexing, Filtering, and Querying Large DataFrames
- Efficient Data Selection and Manipulation

7. Data Cleaning and Preprocessing

- Handling Missing, Duplicate, and Inconsistent Data
- Data Transformation: Apply, Map, and Vectorized Functions
- Data Encoding, Normalization, and Type Optimization
- Working with Time Series and Date-Time Indexes

8. Advanced Data Wrangling

- GroupBy Operations, Aggregations, and Custom Functions
- Merging, Joining, Concatenating, and Reshaping Data
- Pivot Tables, Cross-Tabulations, and Complex Reshaping
- Handling Nested and Hierarchical Data

9. Scaling Pandas for Big Data

- Working with Large Datasets (Out-of-Core Processing)
- Optimizing Memory Usage and Performance
- Using Dask and Vaex for Parallel and Distributed DataFrames
- Chunking and Batch Processing Large Data Files

Part IV: Data Visualization with Matplotlib and Beyond

10. Fundamentals of Data Visualization

- Anatomy of a Matplotlib Plot
- Creating Line Plots, Bar Charts, Histograms, and Scatter Plots
- Customizing Visuals: Colors, Labels, Grids, and Annotations
- Subplots, Multiple Axes, and Plot Styling

11. Advanced Visualization Techniques

- Seaborn for Statistical Visualizations
- Plotly for Interactive and Web-Based Visualizations
- Heatmaps, Correlation Matrices, and Geospatial Plots
- o 3D Plots, Network Graphs, and Complex Visual Representations

12. Building Interactive Dashboards

- Integrating Pandas, Matplotlib, and Plotly for Dashboards
- Using ipywidgets and Voila for Interactivity
- Designing Reproducible Reports with Jupyter and nbconvert

Part V: Advanced Techniques for Experts

13. Optimizing Data Pipelines

- Profiling and Benchmarking Data Operations
- Performance Tuning with Vectorization and Lazy Evaluation
- Handling Sparse Data and Optimizing Data Types
- Memory Mapping and Efficient Disk I/O

14. Memory Management and Scalability

- Reducing Memory Footprint in Pandas and Numpy
- Working with Large-Scale Datasets Using Dask and PySpark
- Parallel and Distributed Data Analysis Techniques
- o Caching, Streaming Data, and Real-Time Analytics

15. Advanced Data Wrangling and Transformation

- Complex Reshaping Techniques (Melt, Pivot, Stack, Unstack)
- Advanced GroupBy, Rolling Windows, and Window Functions
- Time Series Analysis and Anomaly Detection
- Handling High-Dimensional and Multivariate Data

16. Dimensionality Reduction and Complex Data Analysis

- Principal Component Analysis (PCA) and Singular Value Decomposition (SVD)
- o t-SNE, UMAP, and Other Visualization Techniques
- Feature Engineering for Predictive Analytics
- Correlation Analysis and Multicollinearity

Part VI: Real-World Data Analysis Projects

17. Exploratory Data Analysis (EDA) in Practice

- Designing Effective EDA Workflows
- Identifying Patterns, Trends, and Anomalies
- Automating EDA with Custom Scripts
- Case Study: Analyzing E-Commerce Customer Behavior

18. Domain-Specific Data Analysis

- Financial Data Analysis: Time Series, Forecasting, and Risk Modeling
- Healthcare Analytics: Patient Data, Statistics, and Predictive Modeling
- Social Media and Text Data Analytics: Sentiment Analysis with Numpy & Pandas
- Scientific Computing: Monte Carlo Simulations and Experimental Data Analysis

19. Building Scalable Data Analysis Pipelines

- Structuring Complex Data Projects for Reusability
- Data Versioning and Reproducibility with DVC and Git
- Creating Automated ETL Pipelines
- Case Study: End-to-End Churn Prediction Pipeline

Part VII: Expanding Your Toolkit

20. Integrating Data Analysis with External Systems

- Working with Databases (SQL, NoSQL) and APIs
- o Cloud Data Processing with AWS, BigQuery, and Snowflake
- Real-Time Data Streaming with Kafka and Pandas
- Data Engineering for Data Science Workflows

21. From Data Analysis to Data Science

- Preparing Data for Machine Learning Models
- o Introduction to Scikit-learn for Predictive Analytics
- Using Pandas and Numpy for Feature Engineering
- Building and Evaluating Simple Machine Learning Pipelines

22. Best Practices and Next Steps

- Writing Clean, Efficient, and Scalable Code
- Debugging, Profiling, and Memory Optimization
- Ensuring Data Quality and Consistency
- Building Your Data Analysis Portfolio

Appendices

- A. Numpy, Pandas, and Matplotlib Cheat Sheets
- B. Performance Optimization Tips for Data Analysts
- C. Common Data Analysis Pitfalls and How to Avoid Them
- D. Sample Datasets for Practice and Case Studies
- E. Resources for Further Learning (Courses, Books, and Tools)

Chapter 1: Introduction to Data Analysis for **Python Experts**

1. What is Data Analysis?

Definition

Data analysis is the process of inspecting, cleansing, transforming, and modeling data to extract meaningful insights, identify trends, and support decision-making. It bridges the gap between raw data and actionable knowledge.

Key Objectives of Data Analysis:

- Extract insights from large, complex datasets
- Identify patterns, trends, and anomalies
- Support data-driven decision-making
- Predict future outcomes using statistical and machine learning models

III Types of Data Analysis:

Туре	Purpose	Techniques Used
Descriptive Analysis	What happened?	Summarization, Aggregation
Diagnostic Analysis	Why did it happen?	Correlation, Root Cause Analysis
Predictive Analysis	What is likely to happen?	Regression, Classification Models
Prescriptive Analysis	What should we do?	Optimization, Decision Trees
Exploratory Data Analysis (EDA)	Discover hidden patterns	Visualizations, Clustering

2. The Role of Python in Modern Data Analysis

Why Python?

Python has become the leading language for data analysis due to its simplicity, extensive ecosystem, and performance optimization techniques.

Key Strengths:

- **Readable syntax**: Ideal for rapid prototyping and collaboration.
- Rich ecosystem: Libraries like numpy, pandas, matplotlib, and scikit-learn.
- **Scalability**: From small datasets to big data frameworks (Dask, PySpark).
- **Interoperability**: Easy integration with databases, APIs, and cloud services.
- **Community support**: A vast array of tutorials, forums, and open-source projects.

Sesential Python Libraries for Data Analysis:

Library	Purpose	
Numpy	High-performance numerical computations	
Pandas	Data manipulation and analysis	
Matplotlib	Basic plotting and visualization	
Seaborn	Statistical graphics built on top of Matplotlib	
Scikit-learn	Machine learning and predictive modeling	
Dask	Parallel computing for large datasets	
Plotly	Interactive visualizations	

Solution Programming in Data Analysis:

Python's functional programming features (like map, filter, reduce, and comprehensions) allow for concise, expressive data transformations.

Example: Using map and filter for data transformation:

```
# List of sales transactions
sales = [120, 340, 560, 80, 230, 150]

# Apply 10% discount to sales > 200
discounted_sales = list(map(lambda x: x * 0.9 if x > 200 else x, sales))

# Filter out transactions < 100
high_value_sales = list(filter(lambda x: x >= 100, discounted_sales))

print(high_value_sales)
```

Output:

```
[108.0, 306.0, 504.0, 230, 150]
```

3. Understanding the Data Analysis Workflow

A successful data analysis project follows a systematic workflow to ensure accuracy, scalability, and reproducibility.

The 6-Step Data Analysis Workflow:

1. Define Objectives **@**

Understand the problem and set clear goals.

2. Collect Data 📥

- o Pull data from databases, APIs, or files (CSV, JSON, Excel).
- Ensure data is relevant and high-quality.

3. Clean and Prepare Data 🗳

- Handle missing values, duplicates, and data inconsistencies.
- o Normalize, transform, and encode data for analysis.

4. Explore Data (EDA) 🕓

- Use statistical summaries, visualizations, and correlation analysis.
- o Identify trends, patterns, and anomalies.

5. Analyze and Model [1]

- Apply statistical methods or machine learning algorithms.
- Optimize models for accuracy and performance.

6. Visualize and Communicate Insights 🖂

- Build charts, dashboards, and reports to present findings.
- Translate data insights into actionable recommendations.

Example: Analyzing Sales Data

```
import pandas as pd
import matplotlib.pyplot as plt
# Step 1: Load Data
data = pd.DataFrame({
    'Product': ['A', 'B', 'C', 'D', 'E'],
    'Sales': [250, 340, 560, 480, 150],
    'Returns': [5, 7, 9, 3, 4]
})
# Step 2: Calculate Net Sales
data['Net Sales'] = data['Sales'] - data['Returns']
# Step 3: Visualize
plt.bar(data['Product'], data['Net_Sales'], color='skyblue')
plt.title('Net Sales by Product')
plt.xlabel('Product')
plt.ylabel('Net Sales')
plt.show()
```

Visualization Output: A bar chart showing net sales for each product.

4. Setting Up a High-Performance Data Analysis Environment

Recommended Tools & Libraries:

- Jupyter Notebook/Lab → Interactive coding and visualization
- **VSCode or PyCharm** → Advanced IDEs with data science extensions
- **Anaconda Distribution** → Pre-configured Python data science environment
- Virtual Environments → Use venv or conda for dependency management

Installing Essential Libraries:

pip install numpy pandas matplotlib seaborn scikit-learn jupyter

Or using **Anaconda**:

conda install numpy pandas matplotlib seaborn scikit-learn jupyter

Best Practices for Data Analysis Projects:

- **Use version control** (Git) for reproducibility.
- Modularize code into reusable functions and scripts.
- **Document everything** from data sources to analysis steps.
- Automate repetitive tasks using functional programming and pipelines.
- **Optimize performance** by using vectorized operations and memory-efficient data structures.

A Chapter Summary:

- Data analysis turns raw data into actionable insights.
- Python offers an extensive ecosystem for efficient data analysis.
- A **structured workflow** helps ensure clean, scalable, and reproducible analyses.
- Functional programming techniques can simplify data manipulation tasks.
- A high-performance environment sets the foundation for successful data projects.

Chapter 2: Core Concepts in Data Analysis

1. Understanding Structured, Semi-Structured, and Unstructured Data

fin Structured Data

Structured data follows a clear, defined schema, often stored in tabular formats like databases or spreadsheets. It allows for efficient querying and is ideal for traditional data analysis.

• Examples:

- Relational databases (MySQL, PostgreSQL)
- CSV/Excel files
- DataFrames in Pandas

Example in Pandas:

```
import pandas as pd

# Structured data in a DataFrame
data = pd.DataFrame({
    'Product': ['A', 'B', 'C'],
    'Price': [100, 150, 200],
    'Stock': [30, 20, 15]
})

print(data)
```

Output:

```
Product Price Stock
0 A 100 30
1 B 150 20
2 C 200 15
```

Semi-Structured Data

Semi-structured data doesn't adhere to a strict schema but still contains tags or markers to separate elements.

• Examples:

- o JSON, XML
- NoSQL databases (MongoDB)
- API responses

Example (JSON Parsing):

```
import json

# Sample JSON data
json_data = '''
{
    "employees": [
        {"name": "Alice", "role": "Developer"},
        {"name": "Bob", "role": "Designer"}
    ]
}
'''

# Parse JSON
parsed = json.loads(json_data)

# Access data
for emp in parsed['employees']:
    print(f"{emp['name']} - {emp['role']}")
```

Output:

```
Alice - Developer
Bob - Designer
```

Unstructured Data

Unstructured data lacks a predefined format or organization, making it more challenging to process.

- Examples:
 - o Text files, emails
 - o Images, videos
 - Social media posts

Text Data Example (Word Count):

```
from collections import Counter

text = "Data analysis is fun. Data is powerful."
words = text.lower().replace('.', '').split()
word_count = Counter(words)

print(word_count)
```

```
Counter({'data': 2, 'analysis': 1, 'is': 2, 'fun': 1, 'powerful': 1})
```

2. The Data Analysis Pipeline: From Raw Data to Insights

A successful data analysis project follows a logical pipeline, ensuring data integrity and meaningful outcomes.

6-Step Data Analysis Pipeline:

1. Data Collection 🕹

o Gather data from APIs, databases, web scraping, or CSV files.

2. Data Cleaning 🖋

Handle missing values, remove duplicates, and correct inconsistencies.

3. Data Transformation 🖸

Apply scaling, normalization, encoding, and feature engineering.

4. Exploratory Data Analysis (EDA) [[]

• Discover patterns, correlations, and anomalies using visualizations.

5. Data Modeling & Analysis 🖂

Use statistical methods or machine learning for deeper insights.

6. Presentation of Results

• Share findings using dashboards, reports, or visualizations.

Example Pipeline (Sales Data Analysis):

```
import pandas as pd
import matplotlib.pyplot as plt

# Step 1: Load Data
data = pd.DataFrame({
    'Product': ['A', 'B', 'C', 'D'],
    'Sales': [250, 340, 560, 480],
    'Returns': [5, 7, 9, 3]
})

# Step 2: Clean Data (No missing values here)

# Step 3: Transform Data (Net Sales)
data['Net_Sales'] = data['Sales'] - data['Returns']
```

```
# Step 4: EDA (Visualize Net Sales)
plt.bar(data['Product'], data['Net_Sales'], color='lightgreen')
plt.title('Net Sales by Product')
plt.xlabel('Product')
plt.ylabel('Net Sales')
plt.show()
```

3. Exploratory Data Analysis (EDA) vs. Explanatory Analysis

S Exploratory Data Analysis (EDA):

- Purpose: Discover patterns, spot anomalies, and test hypotheses.
- Methods: Descriptive statistics, visualizations, correlation matrices.
- Outcome: Identify trends and areas for deeper investigation.

Example (EDA with Pandas):

```
import seaborn as sns

# Sample Data
df = sns.load_dataset('tips')

# Correlation heatmap
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```

Explanatory Data Analysis:

- Purpose: Communicate specific insights and findings.
- Methods: Polished visualizations, storytelling with data.
- Outcome: Clear and concise presentation of results to stakeholders.

Key Differences:

Aspect	EDA	Explanatory Analysis
Goal	Explore, discover patterns	Communicate insights
Audience	Analysts, Data Scientists	Stakeholders, Decision-Makers
Tools Used	Histograms, Boxplots, Heatmaps	Dashboards, Infographics
Flexibility	Open-ended	Goal-driven

4. Data Ethics, Integrity, and Reproducibility



- Transparency: Always disclose data sources and methods.
- **Privacy:** Follow data protection laws (e.g., GDPR).
- Fairness: Avoid biases that could impact results.
- Accountability: Ensure responsible data handling.

Ш Data Integrity:

- Data Accuracy: Ensure data reflects real-world conditions.
- Consistency: Maintain uniform formats and standards.
- Completeness: Fill in missing data or acknowledge gaps.

Example: Handling Missing Data in Pandas:

```
import pandas as pd

# Sample DataFrame with missing values

df = pd.DataFrame({
        'Name': ['Alice', 'Bob', 'Charlie'],
        'Score': [85, None, 92]
})

# Fill missing values with the mean

df['Score'].fillna(df['Score'].mean(), inplace=True)

print(df)
```

Output:

```
Name Score

0 Alice 85.0

1 Bob 88.5

2 Charlie 92.0
```

Reproducibility in Data Analysis:

- **Version Control:** Use **Git** to track changes in data and code.
- **Documentation:** Keep clear records of data sources, transformations, and decisions.
- Automated Pipelines: Build repeatable workflows using scripts or notebooks.
- Notebook Best Practices: Use Jupyter with clear markdown cells, code explanations, and outputs.

Example: Reproducible Analysis Workflow

- 1. **Data Collection:** API calls → Save raw data
- 2. **Data Cleaning:** Scripted transformations → Save cleaned dataset
- 3. Modeling & Analysis: Version-controlled notebooks
- 4. **Results Sharing:** Export plots/reports → Share as PDFs or interactive dashboards

A Chapter Summary:

- Structured, semi-structured, and unstructured data require different analysis techniques.
- The Data Analysis Pipeline transforms raw data into insights through cleaning, transformation, and modeling.
- EDA helps discover patterns, while Explanatory Analysis communicates results effectively.
- Upholding data ethics, integrity, and reproducibility ensures trust in data-driven decisions.

☐ Chapter 3: Python Essentials for Data Analysis (Quick Recap)

This chapter provides a concise yet powerful recap of essential Python concepts tailored for data analysis. While you're already a Python expert, this chapter will focus on applying core Python techniques effectively in data analysis workflows.

1. Working with Iterators, Generators, and List Comprehensions

@ Iterators

An **iterator** is an object that enables traversing through a collection, one element at a time, without needing to store the entire collection in memory.

• Key methods:

- <u>__iter__()</u> returns the iterator object.
- o next () returns the next element.

Example: Custom Iterator

```
class Counter:
    def __init__(self, low, high):
        self.current = low
        self.high = high

def __iter__(self):
        return self

def __next__(self):
        if self.current > self.high:
            raise StopIteration
        else:
            self.current += 1
            return self.current - 1

# Using the iterator
for num in Counter(1, 5):
        print(num)
```

```
1
2
3
4
5
```

⇔ Generators

Generators simplify iterators using the yield statement. They're memory-efficient, ideal for large datasets.

Example: Data Stream Generator

```
def data_stream(n):
    for i in range(n):
        yield i * 2

stream = data_stream(5)
for value in stream:
    print(value)
```

Output:

```
0
2
4
6
8
```

$\begin{picture}(100,0) \put(0,0){\line(0,0){100}} \put(0,0){\line(0,0){10$

• Streaming large datasets row by row to avoid memory overload.

Example 1 List Comprehensions

List comprehensions offer a concise way to create lists and apply transformations.

Example: Filtering Data

```
# Filter even numbers from a list
numbers = [1, 2, 3, 4, 5, 6]
evens = [x for x in numbers if x % 2 == 0]
print(evens)
```

```
[2, 4, 6]
```

Tip: Use set comprehensions {} and dict comprehensions {key: value} similarly.

1 2. Lambda Functions, Map-Reduce, and Functional Programming Techniques

6 Lambda Functions

Lambda functions are anonymous, inline functions, perfect for short transformations.

Example: Simple Lambda

```
square = lambda x: x ** 2
print(square(5))
```

Output:

```
25
```

Map, Filter, and Reduce

- map() → Applies a function to each item in a list.
- **filter()** → Filters items based on a condition.
- reduce() → Aggregates values (requires functools).

Example: Apply Map and Filter

```
from functools import reduce

numbers = [1, 2, 3, 4, 5]

# Double each number
doubled = list(map(lambda x: x * 2, numbers))

# Filter even numbers
evens = list(filter(lambda x: x % 2 == 0, doubled))

# Sum all numbers
total = reduce(lambda x, y: x + y, evens)

print(f"Doubled: {doubled}")
print(f"Evens: {evens}")
print(f"Sum of Evens: {total}")
```

```
Doubled: [2, 4, 6, 8, 10]
Evens: [2, 4, 6, 8, 10]
Sum of Evens: 30
```

Functional Programming in Data Analysis

Functional programming techniques simplify data pipelines by applying operations in sequence.

Example: Chaining Operations with map and filter

```
data = range(10)

result = sum(
    map(lambda x: x ** 2,
         filter(lambda x: x % 2 == 0, data))
)

print(result)
```

Explanation:

- 1. Filter even numbers.
- 2. **Square** the filtered numbers.
- 3. **Sum** the squared values.

Output:

```
120
```

3. Error Handling and Debugging for Data Analysis

Working with real-world data means encountering errors — missing values, malformed data, or unexpected types. Efficient error handling is crucial.

⚠ Try-Except Blocks

Use try-except to catch runtime errors gracefully.

Example: Handling Division by Zero

```
def safe_divide(a, b):
    try:
        return a / b
    except ZeroDivisionError:
```

```
return "Cannot divide by zero"

print(safe_divide(10, 2)) # 5.0
print(safe_divide(10, 0)) # Cannot divide by zero
```

Debugging Tips for Data Analysis:

• Use assertions to catch anomalies early:

```
assert df.isnull().sum().sum() == 0, "Data contains missing values!"
```

• Leverage pdb for step-by-step debugging:

```
import pdb
pdb.set_trace()
```

• Use logging instead of print statements for better traceability:

```
import logging
logging.basicConfig(level=logging.INFO)
logging.info("Data loaded successfully")
```

4. Efficient File I/O: Working with CSV, JSON, and SQL

CSV Files

Read and write CSV using Pandas:

```
import pandas as pd

# Read CSV

df = pd.read_csv('data.csv')

# Write CSV

df.to_csv('output.csv', index=False)
```

Ⅲ JSON Files

Load and save JSON data:

```
import json

# Sample data
data = {"name": "Alice", "age": 30, "role": "Analyst"}

# Write to JSON
with open('data.json', 'w') as file:
    json.dump(data, file)

# Read JSON
with open('data.json', 'r') as file:
    loaded_data = json.load(file)

print(loaded_data)
```

SQL Databases

Connect and query using sqlite3:

```
import sqlite3
import pandas as pd
# Create connection
conn = sqlite3.connect('example.db')
# Write DataFrame to SQL
df = pd.DataFrame({
    'Name': ['Alice', 'Bob'],
    'Age': [30, 25]
})
df.to_sql('users', conn, if_exists='replace', index=False)
# Read from SQL
query = "SELECT * FROM users"
result = pd.read_sql(query, conn)
print(result)
# Close connection
conn.close()
```

A Chapter Summary:

- Iterators, Generators, and List Comprehensions streamline data workflows.
- Lambda, Map-Reduce, and Functional Programming simplify complex data transformations.
- Robust error handling ensures data integrity during analysis.
- Efficient **File I/O operations** enable seamless data movement between formats and databases.

Chapter 4: Numpy Essentials

NumPy is the cornerstone of numerical computing in Python. Its efficient array structures and high-performance operations make it a foundational tool for data analysis, scientific computing, and machine learning. This chapter will cover essential NumPy concepts and advanced techniques to help you leverage its full potential.

1. Numpy Arrays: The Foundation of Numerical Computing

Why NumPy Arrays Over Python Lists?

- Memory Efficiency: NumPy arrays use less memory compared to Python lists.
- **Speed:** Operations on arrays are significantly faster due to optimized C backend.
- Convenience: Built-in functions for mathematical operations, broadcasting, and aggregation.

Creating Numpy Arrays

```
import numpy as np
# 1D Array
arr1 = np.array([1, 2, 3, 4])
print("1D Array:", arr1)
# 2D Array (Matrix)
arr2 = np.array([[1, 2], [3, 4]])
print("2D Array:\n", arr2)
# Special Arrays
                             # 3x3 matrix of zeros
zeros = np.zeros((3, 3))
ones = np.ones((2, 4))
                              # 2x4 matrix of ones
                              # 3x3 identity matrix
identity = np.eye(3)
random_arr = np.random.rand(2, 3) # 2x3 random numbers
print("Zeros:\n", zeros)
print("Random Array:\n", random arr)
```

Mathematical Report of Numpy Arrays

```
print("Shape:", arr2.shape) # (2, 2)
print("Size:", arr2.size) # 4
print("Data Type:", arr2.dtype) # int64 or float64
```

2. Vectorization: Replacing Loops with Efficient Operations

Vectorization allows operations on entire arrays without explicit loops, resulting in concise and faster code.

VS For Loop vs. Vectorized Operation

```
# Python loop
data = [1, 2, 3, 4]
squared_loop = [x**2 for x in data]

# NumPy vectorized
arr = np.array(data)
squared_np = arr ** 2

print("Loop Result:", squared_loop)
print("NumPy Result:", squared_np)
```

OPERATION Performance Comparison:

```
import time

data = np.random.rand(1000000)

# Python list
start = time.time()
squared_list = [x**2 for x in data]
print("Python list time:", time.time() - start)

# NumPy array
arr = np.array(data)
start = time.time()
squared_np = arr ** 2
print("NumPy time:", time.time() - start)
```

NumPy will be ~10-50x faster due to optimized C routines.

Common Vectorized Operations

```
# Conditional operations
print(arr[arr > 20]) # [30 40]
```

3. Broadcasting, Indexing, and Slicing Techniques

Broadcasting allows NumPy to perform arithmetic operations on arrays of different shapes.

Rules of Broadcasting:

- 1. Dimensions are compared from right to left.
- 2. A dimension of size 1 can be stretched to match.

Indexing and Slicing

Reversing Arrays:

```
arr = np.arange(10)
print(arr[::-1]) # [9 8 7 6 5 4 3 2 1 0]
```

• Selecting Diagonals:

```
matrix = np.arange(1, 10).reshape(3, 3)
print(np.diag(matrix)) # [1 5 9]
```

4. Statistical Functions, Aggregations, and Basic Linear Algebra

Statistical Functions

```
data = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9])

print("Mean:", np.mean(data))
print("Median:", np.median(data))
print("Standard Deviation:", np.std(data))
print("Sum:", np.sum(data))
print("Min:", np.min(data))
print("Max:", np.max(data))
```

M Aggregations Along Axes

Basic Linear Algebra with NumPy

NumPy provides efficient implementations for common linear algebra operations.

```
from numpy.linalg import inv, eig, det

# Matrix operations
A = np.array([[1, 2], [3, 4]])
```

```
# Transpose
print("Transpose:\n", A.T)

# Determinant
print("Determinant:", det(A))

# Inverse
print("Inverse:\n", inv(A))

# Eigenvalues and Eigenvectors
values, vectors = eig(A)
print("Eigenvalues:", values)
print("Eigenvectors:\n", vectors)
```

Solving Linear Systems

```
# Solve Ax = b
A = np.array([[3, 1], [1, 2]])
b = np.array([9, 8])

x = np.linalg.solve(A, b)
print("Solution x:", x)
```

A Chapter Summary:

- NumPy arrays provide efficient storage and computation for numerical data.
- **Vectorization** eliminates the need for slow Python loops.
- **Broadcasting** enables operations on arrays of different shapes.
- Use **indexing**, **slicing**, **and boolean masks** for flexible data access.
- Leverage statistical functions and linear algebra operations for complex analysis.

Chapter 5: Advanced Numpy Techniques

In this chapter, we'll go beyond the basics of NumPy and dive into advanced topics that unlock the full power of numerical computing in Python. We will explore memory management strategies, structured arrays, optimization techniques using **Numba** and **Cython**, and how to perform simulations and complex mathematical operations.

1. Memory Management and Data Alignment

Understanding NumPy's Memory Model

NumPy arrays are stored in contiguous blocks of memory, which enables fast access and manipulation. Efficient memory usage is crucial for handling large datasets.

III Key Concepts:

- **Contiguity:** NumPy arrays can be *C-contiguous* (row-major) or *F-contiguous* (column-major).
- Views vs. Copies:
 - **Views** share memory with the original array (faster, less memory).
 - Copies allocate new memory (safe for independent modifications).

Example: View vs. Copy

```
import numpy as np

arr = np.array([1, 2, 3, 4, 5])
view = arr[1:4]  # This is a view
view[0] = 99
print("Original Array:", arr)  # [1 99 3 4 5]

copy = arr[1:4].copy()
copy[0] = 42
print("After Copy Modification:", arr)  # [1 99 3 4 5]
```

Optimizing Memory Usage

• Data Types (dtype) Optimization:

```
# Using smaller data types
arr_float64 = np.arange(1e6, dtype=np.float64)
arr_float32 = np.arange(1e6, dtype=np.float32)

print("Float64 size:", arr_float64.nbytes, "bytes")
print("Float32 size:", arr_float32.nbytes, "bytes")
```

• **Using np. memmap for Large Datasets:** Memory-map large files to avoid loading everything into RAM.

```
data = np.memmap('large_array.dat', dtype='float32', mode='w+', shape=(10000,
10000))
data[5000, 5000] = 42
data.flush() # Ensure data is written to disk
```

2. Structured Arrays and Record Arrays

Structured arrays allow you to store complex datasets with mixed data types — similar to SQL tables or Pandas DataFrames but with lower-level control.

Creating Structured Arrays:

```
# Define a structured dtype
dt = np.dtype([('Name', 'U10'), ('Age', 'i4'), ('Salary', 'f4')])

# Create the array
data = np.array([('Alice', 25, 50000.0), ('Bob', 30, 60000.0)], dtype=dt)

print(data)
```

Output:

```
[('Alice', 25, 50000.) ('Bob', 30, 60000.)]
```

Accessing Structured Data:

```
# Access specific fields
print(data['Name']) # ['Alice' 'Bob']
print(data['Salary']) # [50000. 60000.]

# Conditional filtering
print(data[data['Age'] > 26]) # [('Bob', 30, 60000.)]
```

Record Arrays (recarray) for Attribute Access:

```
rec_data = data.view(np.recarray)
print(rec_data.Name) # ['Alice' 'Bob']
```

When to use:

- Use structured arrays for lightweight data manipulation.
- Use Pandas for heavy-duty data analysis and complex indexing.

3. Optimizing Computations with Numba and Cython

When pure NumPy isn't fast enough, tools like **Numba** and **Cython** can supercharge performance.



Installation:

```
pip install numba
```

Example: Speeding Up Loops

```
from numba import jit
import numpy as np
import time
@jit(nopython=True)
def slow_sum(arr):
   total = 0
    for x in arr:
        total += x
    return total
arr = np.random.rand(10**7)
# Without Numba
start = time.time()
print(np.sum(arr))
print("Numpy Time:", time.time() - start)
# With Numba
start = time.time()
print(slow sum(arr))
print("Numba Time:", time.time() - start)
```

Result: Numba often achieves **5-50x** speedups over regular Python loops.

2 Cython: C Extensions for Python

Installation:

```
pip install cython
```

Example: Accelerating Python with Cython

1. Create a cython_example.pyx file:

```
def cython_sum(double[:] arr):
    cdef double total = 0
    for i in range(arr.shape[0]):
        total += arr[i]
    return total
```

2. Compile with:

```
cythonize -i cython_example.pyx
```

3. Use in Python:

```
import numpy as np
from cython_example import cython_sum

arr = np.random.rand(10**7)
print(cython_sum(arr))
```

Cython is great for loops that can't be vectorized easily.

4. Random Sampling, Simulations, and Complex Mathematical Operations

Random Sampling with numpy.random

```
rng = np.random.default_rng(seed=42)

# Random integers
ints = rng.integers(low=0, high=10, size=5)
print("Random Integers:", ints)

# Random floats
floats = rng.random(size=5)
print("Random Floats:", floats)

# Normal distribution
```

```
normal = rng.normal(loc=0, scale=1, size=5)
print("Normal Distribution:", normal)
```

Monte Carlo Simulation Example

Estimate π using random sampling:

```
import numpy as np

def estimate_pi(n_samples=1000000):
    x = np.random.uniform(0, 1, n_samples)
    y = np.random.uniform(0, 1, n_samples)
    inside_circle = (x**2 + y**2) <= 1
    pi_estimate = 4 * np.sum(inside_circle) / n_samples
    return pi_estimate

print("Estimated π:", estimate_pi())</pre>
```

Complex Mathematical Operations

NumPy supports advanced math operations directly:

• Fourier Transforms:

```
from numpy.fft import fft

x = np.random.rand(100)
y = fft(x)
print("FFT Result:", y)
```

• Polynomials:

```
p = np.poly1d([1, 2, 1]) # x^2 + 2x + 1
print(p(3)) # Evaluate at x=3
```

• Matrix Factorization (SVD):

```
A = np.random.rand(3, 3)
U, S, Vt = np.linalg.svd(A)
print("Singular Values:", S)
```

A Chapter Summary:

- Memory management helps optimize large datasets.
- Structured arrays provide flexible data storage for mixed types.
- Numba and Cython accelerate computations for heavy-duty tasks.
- Random sampling and simulations enable probabilistic modeling and statistical analysis.

🛄 Chapter 6: Pandas Fundamentals

Pandas is the go-to library for data manipulation and analysis in Python. Its core data structures—Series, DataFrames, and MultiIndex—make working with structured data simple and intuitive. In this chapter, we'll cover the fundamental components of Pandas and how to efficiently import, query, and manipulate large datasets.

3. 1. Series, DataFrames, and MultiIndex Structures

1.1 Series: The One-Dimensional Labeled Array

A **Series** is a one-dimensional array with labels (indexes).

```
import pandas as pd
# Creating a Series
data = pd.Series([10, 20, 30, 40], index=['a', 'b', 'c', 'd'])
print(data)
# Accessing by label
print("Element 'b':", data['b'])
# Vectorized operations
print("Doubled values:\n", data * 2)
```

Output:

```
10
     20
     30
dtype: int64
Element 'b': 20
Doubled values:
     20
     40
     60
d
     80
dtype: int64
```

Tip: Series can be treated like both a NumPy array (for math ops) and a dictionary (for key-based access).

1.2 DataFrames: The 2D Data Powerhouse

A **DataFrame** is a 2D labeled data structure with rows and columns.

Output:

```
Name Age Salary

0 Alice 25 70000

1 Bob 30 80000

2 Charlie 35 90000
```

Key DataFrame Operations:

```
# Access columns
print(df['Name'])

# Add a new column
df['Bonus'] = df['Salary'] * 0.10

# Basic statistics
print(df.describe())

# Renaming columns
df.rename(columns={'Salary': 'Annual Salary'}, inplace=True)
print(df)
```

1.3 MultiIndex: Working with Hierarchical Data

MultiIndex enables multi-level (hierarchical) indexing in DataFrames.

```
# Creating a MultiIndex DataFrame
arrays = [['USA', 'USA', 'Canada', 'Canada'], ['NY', 'CA', 'ON', 'QC']]
index = pd.MultiIndex.from_arrays(arrays, names=('Country', 'State'))

data = pd.DataFrame({'Population': [20, 40, 10, 15]}, index=index)
print(data)
```

```
Population

Country State

USA NY 20

CA 40

Canada ON 10

QC 15
```

Accessing with MultiIndex:

```
print(data.loc['USA'])
```

2. Data Import/Export (CSV, Excel, JSON, SQL, APIs)

Pandas provides powerful I/O tools for reading/writing data from various sources.

2.1 CSV Files

```
# Read CSV

df = pd.read_csv('data.csv')

# Write to CSV

df.to_csv('output.csv', index=False)
```

11 2.2 Excel Files

```
# Read Excel
df = pd.read_excel('data.xlsx', sheet_name='Sheet1')

# Write to Excel
df.to_excel('output.xlsx', index=False)
```

3.3 JSON Files

```
# Read JSON
df = pd.read_json('data.json')

# Convert DataFrame to JSON
df.to_json('output.json', orient='records')
```

3.4 SQL Databases

```
import sqlite3

# Connect to SQLite database
conn = sqlite3.connect('my_database.db')

# Read SQL table
df = pd.read_sql('SELECT * FROM employees', conn)

# Write DataFrame to SQL
df.to_sql('new_table', conn, if_exists='replace', index=False)
```

🕸 2.5 APIs (Using JSON)

```
import requests
response = requests.get('https://jsonplaceholder.typicode.com/posts')
json_data = response.json()

# Convert API data to DataFrame
df = pd.DataFrame(json_data)
print(df.head())
```

3. Indexing, Filtering, and Querying Large DataFrames

☆ 3.1 Indexing Techniques

```
# Setting index
df.set_index('Name', inplace=True)

# Resetting index
df.reset_index(inplace=True)
```

3.2 Filtering Data

```
# Boolean indexing
filtered_df = df[df['Age'] > 30]

# Multiple conditions
filtered_df = df[(df['Age'] > 25) & (df['Salary'] > 75000)]
```

♀ 3.3 Efficient Querying with .query()

```
# Using query for better readability
high_salary = df.query('Salary > 75000')
print(high_salary)
```

3.4 Handling Large Datasets

• Load data in chunks:

```
chunks = pd.read_csv('large_data.csv', chunksize=100000)
for chunk in chunks:
   print(chunk.shape)
```

• Optimize memory usage:

```
df = pd.read_csv('large_data.csv', dtype={'ID': 'int32', 'Salary':
   'float32'})
print(df.info())
```

% 4. Efficient Data Selection and Manipulation

- 4.1 Selecting Data with .loc[] and .iloc[]
 - loc[] → label-based selection
 - iloc[] → position-based selection

```
# loc example
print(df.loc[0, 'Name'])

# iloc example
print(df.iloc[1, 2]) # Row 1, Column 2
```

% 4.2 Data Transformation

• Apply functions to columns:

```
df['Age Plus 5'] = df['Age'].apply(lambda x: x + 5)
```

• Vectorized string operations:

```
df['Name'] = df['Name'].str.upper()
```

• Handling missing data:

```
df.fillna(0, inplace=True)  # Replace NaN with 0
df.dropna(inplace=True)  # Drop rows with NaN
```

4.3 Grouping and Aggregating Data

```
# Group by and aggregate
grouped = df.groupby('Department')['Salary'].agg(['mean', 'max', 'min'])
print(grouped)
```

4.4 Pivot Tables and Cross Tabulations

```
# Pivot Table
pivot = pd.pivot_table(df, values='Salary', index='Department', columns='Gender',
aggfunc='mean')

# Cross Tabulation
cross_tab = pd.crosstab(df['Department'], df['Gender'])
```

A Chapter Summary:

- Series and DataFrames form the foundation of Pandas data structures.
- Efficiently **import/export** data from various sources (CSV, Excel, JSON, SQL, APIs).
- Mastered **indexing**, **filtering**, **and querying** for small and large datasets.
- Learned **grouping, aggregation**, and advanced data manipulation techniques.

Chapter 7: Data Cleaning and Preprocessing

Data cleaning and preprocessing is a crucial step in any data analysis pipeline. Dirty data can lead to inaccurate insights and misleading conclusions. In this chapter, we'll cover strategies to handle missing, duplicate, and inconsistent data, apply data transformations, encode and normalize data, optimize data types, and work effectively with time series data.

1. Handling Missing, Duplicate, and Inconsistent Data

? 1.1 Identifying Missing Data

Pandas uses NaN to represent missing data.

Output:

```
Name Age Salary

0 False False False
1 False False False
2 False True False
3 False False True
4 True False False

Name 1
Age 1
Salary 1
dtype: int64
```

% 1.2 Handling Missing Data

• Drop missing values:

```
# Drop rows with any missing values
df_drop = df.dropna()

# Drop columns with missing values
df_drop_col = df.dropna(axis=1)
```

• Fill missing values:

```
# Fill with a constant
df_fill = df.fillna(0)

# Forward fill (propagate last valid observation)
df_ffill = df.fillna(method='ffill')

# Backward fill
df_bfill = df.fillna(method='bfill')

# Fill with mean
df['Age'].fillna(df['Age'].mean(), inplace=True)
```

☐ 1.3 Handling Duplicates

```
# Sample DataFrame with duplicates

df_dup = pd.DataFrame({
        'Name': ['Alice', 'Bob', 'Alice', 'David', 'Bob'],
        'Age': [25, 30, 25, 40, 30]
})

# Detect duplicates
print(df_dup.duplicated())

# Drop duplicate rows
df_unique = df_dup.drop_duplicates()
```

1.4 Resolving Inconsistent Data

Fix inconsistent casing:

```
df['Name'] = df['Name'].str.title() # 'alice' -> 'Alice'
```

Remove extra spaces:

```
df['Name'] = df['Name'].str.strip()
```

• Replace inconsistent entries:

```
df['Name'].replace({'Alicia': 'Alice'}, inplace=True)
```

2. Data Transformation: Apply, Map, and Vectorized Functions

2.1 Vectorized Operations

Pandas leverages NumPy under the hood, enabling vectorized operations for speed.

```
df['Salary After Tax'] = df['Salary'] * 0.7
```

👰 2.2 Map, Apply, and Applymap

- map() → element-wise operation for Series
- apply() → apply function across axis (rows/columns)
- applymap() → element-wise operation for entire DataFrame

```
# Using map on a Series
df['Age Category'] = df['Age'].map(lambda x: 'Young' if x < 30 else 'Old')

# Using apply on DataFrame
df['Tax'] = df['Salary'].apply(lambda x: x * 0.3 if pd.notnull(x) else 0)

# applymap for entire DataFrame
df_numeric = df[['Age', 'Salary']]
df_scaled = df_numeric.applymap(lambda x: x / 1000)</pre>
```

2.3 Complex Transformations with pipe()

pipe() enables functional-style chaining for cleaner code.

```
def add_bonus(df, bonus):
    df['Salary'] += bonus
    return df

df = df.pipe(add_bonus, bonus=5000)
```

3. Data Encoding, Normalization, and Type Optimization

3.1 Encoding Categorical Variables

One-Hot Encoding:

```
df = pd.DataFrame({'Gender': ['Male', 'Female', 'Male']})
df_encoded = pd.get_dummies(df, columns=['Gender'], drop_first=True)
```

• Label Encoding:

```
df['Gender_Code'] = df['Gender'].map({'Male': 1, 'Female': 0})
```

📏 3.2 Normalization and Scaling

• Min-Max Scaling:

```
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
df[['Salary']] = scaler.fit_transform(df[['Salary']])
```

• Standardization (Z-score):

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
df[['Age']] = scaler.fit_transform(df[['Age']])
```

3.3 Optimizing Data Types for Memory Efficiency

```
# Check memory usage
print(df.info(memory_usage='deep'))

# Convert float64 to float32
df['Salary'] = df['Salary'].astype('float32')

# Convert object to category
df['Gender'] = df['Gender'].astype('category')
```

Tip: Optimizing data types can significantly reduce memory usage when working with large datasets.

31 4. Working with Time Series and Date-Time Indexes

4.1 Parsing and Converting Dates

```
# Convert to datetime
df['Join Date'] = pd.to_datetime(df['Join Date'])

# Extract date parts
df['Year'] = df['Join Date'].dt.year
df['Month'] = df['Join Date'].dt.month
```

4.2 Setting Date-Time as Index

```
# Set date as index for time series analysis
df.set_index('Join Date', inplace=True)

# Resample to monthly data
df_monthly = df.resample('M').mean()
```

4.3 Time Series Operations

• Date Range Creation:

```
# Create a range of dates
date_rng = pd.date_range(start='2023-01-01', end='2023-12-31', freq='M')
```

Rolling Window Analysis:

```
# 3-month rolling average
df['Rolling_Mean'] = df['Salary'].rolling(window=3).mean()
```

• Shifting Data:

```
# Lagging by 1 period
df['Salary_Lag1'] = df['Salary'].shift(1)
```

A Chapter Summary:

- Handled missing, duplicate, and inconsistent data using Pandas functions.
- Applied data transformations using map(), apply(), and vectorized operations.
- Performed data encoding, normalization, and optimized data types for efficiency.
- Managed **time series data** with datetime indexes and rolling statistics.

Chapter 8: Advanced Data Wrangling

Data wrangling is the process of transforming raw data into a structured and clean format for analysis. In this chapter, we'll dive deep into **grouping**, **aggregating**, **merging**, **reshaping**, and **handling complex data structures** using Pandas. You'll learn how to efficiently manipulate large datasets, enabling you to extract meaningful insights.

1. GroupBy Operations, Aggregations, and Custom Functions

The groupby() method in Pandas is a powerful tool for splitting data into groups based on specific criteria, performing operations on each group, and then combining the results.

1.1 Basic GroupBy and Aggregation

```
import pandas as pd

# Sample DataFrame
df = pd.DataFrame({
    'Department': ['HR', 'HR', 'IT', 'IT', 'Finance', 'Finance'],
    'Employee': ['Alice', 'Bob', 'Charlie', 'David', 'Eva', 'Frank'],
    'Salary': [70000, 80000, 900000, 85000, 75000, 720000],
    'Bonus': [5000, 6000, 80000, 75000, 4000, 45000]
})

# Group by Department and calculate mean salary
grouped = df.groupby('Department')['Salary'].mean()
print(grouped)
```

Output:

```
Department
Finance 73500.0
HR 75000.0
IT 87500.0
Name: Salary, dtype: float64
```

1.2 Multiple Aggregations with agg()

```
# Apply multiple aggregations
agg_result = df.groupby('Department').agg({
    'Salary': ['mean', 'max', 'min'],
    'Bonus': 'sum'
```

```
})
print(agg_result)
```

Output:

```
Salary
                                     Bonus
               mean
                             min
                                    sum
                      max
Department
Finance
            73500.0 75000 72000
                                    8500
HR
            75000.0 80000 70000
                                   11000
            87500.0 90000 85000
IT
                                   15500
```

% 1.3 Custom Aggregations with apply()

You can use apply() to run custom functions on grouped data.

```
# Custom function to calculate range
def salary_range(x):
    return x.max() - x.min()

# Apply custom function
df_grouped = df.groupby('Department')['Salary'].apply(salary_range)
print(df_grouped)
```

Output:

```
Department
Finance 3000
HR 10000
IT 5000
Name: Salary, dtype: int64
```

1.4 Transforming Data with transform()

• transform() returns a DataFrame of the same shape as the input, allowing for broadcasting operations.

```
# Normalize salary within each department

df['Normalized Salary'] = df.groupby('Department')['Salary'].transform(lambda x:
   (x - x.mean()) / x.std())
   print(df)
```

2. Merging, Joining, Concatenating, and Reshaping Data

Combining datasets is essential for data wrangling. Pandas offers flexible functions for merging and joining datasets.


```
# Sample DataFrames
df_employees = pd.DataFrame({
    'Employee': ['Alice', 'Bob', 'Charlie', 'David'],
    'Department': ['HR', 'HR', 'IT', 'IT']
})
df_salaries = pd.DataFrame({
    'Employee': ['Alice', 'Bob', 'Charlie', 'David'],
    'Salary': [70000, 80000, 90000, 85000]
})
# Merge on 'Employee'
merged_df = pd.merge(df_employees, df_salaries, on='Employee')
print(merged_df)
```

Output:

```
Employee Department Salary
   Alice
                 HR
                     70000
1
      Bob
                 HR
                      80000
2 Charlie
                 ΙT
                      90000
3
  David
                 ΙT
                    85000
```

9 2.2 Types of Joins

- Inner Join → Keeps only matching rows.
- **Left Join** → Keeps all rows from the left DataFrame.
- **Right Join** → Keeps all rows from the right DataFrame.
- Outer Join → Keeps all rows from both DataFrames.

```
# Outer join example
df1 = pd.DataFrame({'ID': [1, 2, 3], 'Name': ['Alice', 'Bob', 'Charlie']})
df2 = pd.DataFrame({'ID': [2, 3, 4], 'Salary': [80000, 90000, 85000]})
outer_join = pd.merge(df1, df2, on='ID', how='outer')
print(outer_join)
```

Output:

```
ID Name Salary
0 1 Alice NaN
1 2 Bob 80000.0
2 3 Charlie 90000.0
3 4 NaN 85000.0
```

2.3 Concatenating DataFrames with concat()

```
# Concatenate vertically
df1 = pd.DataFrame({'Name': ['Alice', 'Bob'], 'Age': [25, 30]})
df2 = pd.DataFrame({'Name': ['Charlie', 'David'], 'Age': [35, 40]})

df_concat = pd.concat([df1, df2], ignore_index=True)
print(df_concat)
```

Output:

```
Name Age

0 Alice 25

1 Bob 30

2 Charlie 35

3 David 40
```

2.4 Reshaping Data with melt() and pivot()

- melt() → Unpivot DataFrame from wide to long format.
- pivot() → Convert long format back to wide.

Output:

```
Name Subject Score

0 Alice Math 85

1 Bob Math 90

2 Alice English 78

3 Bob English 82
```

III 3. Pivot Tables, Cross-Tabulations, and Complex Reshaping

3.1 Pivot Tables for Data Summarization

```
df = pd.DataFrame({
    'Department': ['HR', 'HR', 'IT', 'IT', 'Finance', 'Finance'],
    'Gender': ['Female', 'Male', 'Female', 'Female', 'Male'],
    'Salary': [70000, 80000, 90000, 85000, 75000, 72000]
})

# Create Pivot Table
pivot_table = pd.pivot_table(df, values='Salary', index='Department',
columns='Gender', aggfunc='mean')
print(pivot_table)
```

Output:

```
Gender Female Male
Department
Finance 75000.0 72000.0
HR 70000.0 80000.0
IT 85000.0 90000.0
```

3.2 Cross-Tabulations with crosstab()

Useful for frequency tables.

```
# Crosstab example
crosstab = pd.crosstab(df['Department'], df['Gender'])
print(crosstab)
```

Output:

```
Gender Female Male
Department
Finance 1 1
```

```
HR 1 1
IT 1 1
```

♦ 3.3 MultiIndex Reshaping with stack() and unstack()

```
# Create MultiIndex DataFrame
df_multi = df.set_index(['Department', 'Gender'])

# Stack the DataFrame
stacked = df_multi.stack()

# Unstack to reshape
unstacked = df_multi.unstack()
```

4. Handling Nested and Hierarchical Data

4.1 Working with Nested JSON Data

Output:

```
name details_age details_city
0 Alice 25 NY
1 Bob 30 LA
```

4.2 Flattening Hierarchical Data

```
# Multi-level column DataFrame
df = pd.DataFrame({
      ('Employee', 'Name'): ['Alice', 'Bob'],
      ('Employee', 'Age'): [25, 30],
      ('Salary', 'Base'): [70000, 80000],
      ('Salary', 'Bonus'): [5000, 6000]
})

# Flatten columns
df.columns = ['_'.join(col) for col in df.columns]
print(df)
```

Output:

```
Employee_Name Employee_Age Salary_Base Salary_Bonus
0 Alice 25 70000 5000
1 Bob 30 80000 6000
```

A Chapter Summary:

- Mastered grouping, aggregating, and applying custom functions.
- Efficiently **merged**, **joined**, and **concatenated** multiple datasets.
- Built complex pivot tables, cross-tabs, and used reshaping tools.
- Flattened and worked with **nested JSON** and **hierarchical data** structures.

oxdot Chapter 9: Scaling Pandas for Big Data

Pandas is powerful for data manipulation but struggles with extremely large datasets that exceed system memory. In this chapter, we'll explore strategies to scale Pandas for big data workflows, including memory optimizations, out-of-core processing, and leveraging libraries like Dask and Vaex for parallel and distributed computing.

1. Working with Large Datasets (Out-of-Core Processing)

Out-of-core processing allows you to work with datasets that don't fit into memory by loading data in chunks.

1.1 Reading Data in Chunks

Use the chunksize parameter in Pandas' read_csv() to process large files incrementally.

```
import pandas as pd
# Load CSV in chunks of 100,000 rows
chunksize = 100 000
chunk_iter = pd.read_csv('large_dataset.csv', chunksize=chunksize)
# Process chunks
for chunk in chunk_iter:
    chunk_result = chunk.groupby('Category')['Sales'].sum()
    print(chunk_result.head())
```

This reads and processes the file in manageable pieces rather than loading it all at once.

IIII 1.2 Incremental Aggregation Example

If you want to compute global statistics across chunks:

```
# Initialize empty DataFrame
final result = pd.Series(dtype='float64')
# Aggregate across chunks
for chunk in pd.read_csv('large_dataset.csv', chunksize=chunksize):
    chunk_result = chunk.groupby('Region')['Sales'].sum()
    final_result = final_result.add(chunk_result, fill_value=0)
print(final result)
```

② 2. Optimizing Memory Usage and Performance

Handling big data efficiently requires reducing memory consumption.

2.1 Reducing Memory Footprint

2.1.1 Check Memory Usage

```
df = pd.read_csv('large_dataset.csv')
print(df.info(memory_usage='deep'))
```

2.1.2 Downcasting Data Types

Convert larger data types to more efficient ones where possible.

```
# Convert float64 to float32 and int64 to int32
df['Sales'] = pd.to_numeric(df['Sales'], downcast='float')
df['Quantity'] = pd.to_numeric(df['Quantity'], downcast='integer')
```

2.1.3 Convert Object Columns to Category

If a column has a limited number of unique values, convert it to category to save space.

```
df['Product'] = df['Product'].astype('category')
```

 $\begin{picture}(100,0)\put(0,0){\line(0,0){100}}\put(0,0$

2.2 Efficient Data Loading

2.2.1 Selective Column Loading

Only load necessary columns using the usecols parameter.

```
df = pd.read_csv('large_dataset.csv', usecols=['Date', 'Product', 'Sales'])
```

2.2.2 Parse Dates While Reading

Parsing dates during read avoids post-processing.

```
df = pd.read_csv('large_dataset.csv', parse_dates=['OrderDate'])
```


When Pandas is not enough, use **Dask** or **Vaex** for parallel processing and out-of-core data analysis.

3.1 Dask: Parallel Computing with Pandas Syntax

Dask offers a parallel DataFrame similar to Pandas but optimized for large datasets.

3.1.1 Dask Setup and Basics

```
import dask.dataframe as dd

# Load large CSV with Dask
dask_df = dd.read_csv('large_dataset.csv')

# Perform aggregation
result = dask_df.groupby('Region')['Sales'].mean().compute()
print(result)
```

- Lazy Evaluation: Operations are queued until .compute() is called.
- Parallelization: Dask splits the data into partitions and processes them in parallel.

3.1.2 Optimizing Dask Workflows

- Use persist() to cache data in memory across computations.
- Monitor performance using Dask's dashboard (dask.distributed).

```
from dask.distributed import Client

client = Client() # Starts a local Dask cluster
```

♦ 3.2 Vaex: Fast, Out-of-Core DataFrames

Vaex handles **billions of rows** efficiently without loading data into memory.

3.2.1 Vaex Setup and Usage

```
import vaex
```

```
# Load large CSV with Vaex
vaex_df = vaex.from_csv('large_dataset.csv', convert=True)

# Perform lazy evaluation
result = vaex_df.groupby('Region', agg=vaex.agg.mean('Sales'))
print(result)
```

- Zero-Copy Memory Mapping: Fast file access without loading the entire dataset.
- Lazy Evaluation: Only computes when needed.

4. Chunking and Batch Processing Large Data Files

For massive datasets stored in flat files or databases, processing data in **batches** ensures stability.

4.1 Writing Large DataFrames in Chunks

When saving large data, write in chunks to avoid memory issues.

```
# Write large DataFrame in chunks
for i, chunk in enumerate(pd.read_csv('large_dataset.csv', chunksize=50000)):
    chunk.to_csv(f'chunk_{i}.csv', index=False)
```

4.2 Batch Processing with SQL Databases

When data is stored in databases, read and process data in batches using **SQLAlchemy**.

```
from sqlalchemy import create_engine

# Connect to database
engine = create_engine('sqlite:///sales_data.db')

# Read in batches
query = 'SELECT * FROM sales_data'
chunks = pd.read_sql(query, engine, chunksize=100000)

for chunk in chunks:
    print(chunk['Sales'].sum())
```

4.3 Handling Compressed and Remote Data

Pandas can directly read compressed and remote files.

```
# Reading compressed CSV
df = pd.read_csv('data.csv.gz', compression='gzip')

# Reading data from a URL
df = pd.read_csv('https://example.com/large_dataset.csv')
```

A Chapter Summary

✓ Processed large datasets using **out-of-core** techniques. ✓ Optimized memory usage with **downcasting** and **categorical data types**. ✓ Leveraged **Dask** and **Vaex** for parallel and distributed computations. ✓ Used **chunking** and **batch processing** for working with large files.

Chapter 10: Fundamentals of Data Visualization

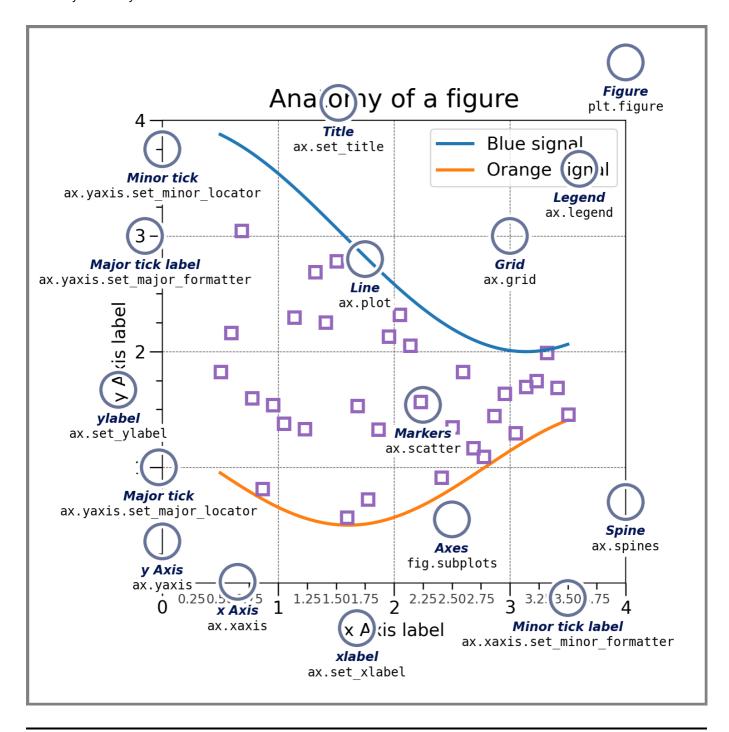
Data visualization is a crucial part of data analysis, enabling you to interpret complex datasets through clear and impactful graphics. In this chapter, we'll dive into **Matplotlib**, the foundational plotting library in Python, covering its core concepts, essential plots, and powerful customization options.

1. Anatomy of a Matplotlib Plot

Before diving into creating visualizations, it's important to understand the structure of a Matplotlib plot.

\(\) 1.1 Core Components

- **Figure** → The entire plotting window (container for one or more plots).
- **Axes** → The actual plot area inside the figure (can have multiple axes).
- **Axis** → The x-axis and y-axis of a plot.
- **Ticks** → Markers on the axes indicating specific data points.
- **Labels** → Text annotations for axes and data points.
- **Legend** → Describes data series within the plot.



1.2 Creating a Basic Plot

```
import matplotlib.pyplot as plt
import numpy as np

# Sample Data
x = np.linspace(0, 10, 100)
y = np.sin(x)

# Create Figure and Axes
fig, ax = plt.subplots()

# Plotting
ax.plot(x, y, label='Sine Wave', color='blue')
```

```
# Add Labels and Title
ax.set_title('Sine Wave Plot')
ax.set_xlabel('X-axis')
ax.set_ylabel('Y-axis')
ax.legend()

# Show Plot
plt.show()
```

11 2. Creating Line Plots, Bar Charts, Histograms, and Scatter Plots

Let's dive into the most common types of visualizations used in data analysis.

2.1 Line Plots

Line plots are perfect for showing trends over time or continuous data.

```
x = np.linspace(0, 10, 100)
y1 = np.sin(x)
y2 = np.cos(x)

plt.figure(figsize=(8,5))
plt.plot(x, y1, label='Sine', color='blue', linestyle='--')
plt.plot(x, y2, label='Cosine', color='red')
plt.title('Sine and Cosine Functions')
plt.xlabel('X-axis')
plt.ylabel('Y-axis')
plt.legend()
plt.show()
```


Bar charts are used for comparing categorical data.

```
categories = ['A', 'B', 'C', 'D']
values = [23, 45, 56, 12]

plt.bar(categories, values, color='teal')
plt.title('Category Comparison')
plt.xlabel('Categories')
plt.ylabel('Values')
plt.show()
```

• Use plt.barh() for horizontal bar charts.

2.3 Histograms

Histograms visualize the distribution of numerical data.

```
data = np.random.randn(1000)

plt.hist(data, bins=30, color='coral', edgecolor='black')
plt.title('Data Distribution')
plt.xlabel('Value')
plt.ylabel('Frequency')
plt.show()
```

• **Tip:** Adjust bins to control granularity.

② 2.4 Scatter Plots

Scatter plots are ideal for examining relationships between variables.

```
x = np.random.rand(100)
y = np.random.rand(100)
sizes = 300 * np.random.rand(100) # Varying bubble sizes
colors = np.random.rand(100)

plt.scatter(x, y, s=sizes, c=colors, alpha=0.6, cmap='viridis')
plt.title('Scatter Plot Example')
plt.xlabel('X-axis')
plt.ylabel('Y-axis')
plt.colorbar(label='Color Intensity')
plt.show()
```

• Add a **colorbar** for extra data dimensions.

% 3. Customizing Visuals: Colors, Labels, Grids, and Annotations

Customizations can make your plots clearer and more engaging.

3.1 Colors and Styles

- Use color names ('red'), HEX codes ('#FF5733'), or color maps.
- Apply line styles: '-' (solid), '--' (dashed), ':' (dotted).

```
x = np.linspace(0, 10, 100)
y = np.sin(x)
```

```
plt.plot(x, y, color='#1f77b4', linestyle='--', linewidth=2, marker='o')
plt.title('Customized Line Plot')
plt.show()
```

3.2 Adding Labels and Legends

- set_xlabel() and set_ylabel() for axes labels.
- legend() for data series.

```
plt.plot(x, y, label='Sine Wave')
plt.xlabel('Time')
plt.ylabel('Amplitude')
plt.title('Labeled Plot')
plt.legend(loc='upper right')
plt.show()
```

3.3 Grids and Annotations

- Add grids for better readability: plt.grid(True)
- Use annotations to highlight specific points.

4. Subplots, Multiple Axes, and Plot Styling

For comparative analysis, you can create complex layouts with multiple plots.

4.1 Creating Subplots

Use plt.subplots() to create multi-plot grids.

```
fig, axs = plt.subplots(2, 2, figsize=(10, 8))
x = np.linspace(0, 10, 100)
axs[0, 0].plot(x, np.sin(x), 'r')
axs[0, 0].set_title('Sine')
```

```
axs[0, 1].plot(x, np.cos(x), 'b')
axs[0, 1].set_title('Cosine')

axs[1, 0].plot(x, np.tan(x), 'g')
axs[1, 0].set_title('Tangent')

axs[1, 1].hist(np.random.randn(1000), bins=30, color='purple')
axs[1, 1].set_title('Histogram')

plt.tight_layout()
plt.show()
```

4.2 Twin Axes for Different Scales

Overlay plots with different y-axis scales.

```
fig, ax1 = plt.subplots()

x = np.arange(0, 10, 0.1)
y1 = np.sin(x)
y2 = y1**2

# Primary axis
ax1.plot(x, y1, 'b-')
ax1.set_xlabel('X-axis')
ax1.set_ylabel('Sine', color='b')

# Secondary axis
ax2 = ax1.twinx()
ax2.plot(x, y2, 'r--')
ax2.set_ylabel('Sine Squared', color='r')

plt.title('Twin Axes Example')
plt.show()
```

4.3 Applying Plot Styles

Matplotlib comes with pre-defined styles like ggplot, seaborn, and fivethirtyeight.

```
plt.style.use('ggplot')

x = np.linspace(0, 10, 100)
y = np.sin(x)

plt.plot(x, y)
plt.title('Styled Plot with ggplot Theme')
plt.show()
```

View all styles:

print(plt.style.available)

✓ Chapter Summary

♦ Mastered the anatomy of Matplotlib plots. ♦ Created line plots, bar charts, histograms, and scatter plots. ♦ Applied customizations for enhanced clarity and aesthetics. ♦ Built complex layouts with subplots and multiple axes.

Chapter 11: Advanced Visualization Techniques

In this chapter, we'll explore advanced data visualization tools and techniques that go beyond basic plotting. From **statistical visualizations** using **Seaborn** to **interactive plots** with **Plotly**, and from **correlation heatmaps** to **geospatial** and **3D plots**, these methods will enable you to present complex data more intuitively and effectively.

1. Seaborn for Statistical Visualizations

Seaborn is built on top of Matplotlib, providing a high-level interface for drawing attractive and informative statistical graphics.

♀ 1.1 Seaborn Setup

```
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

# Load Seaborn's built-in dataset
tips = sns.load_dataset('tips')
```

III 1.2 Enhanced Categorical Plots

1.2.1 Barplot with Confidence Intervals

```
sns.barplot(x='day', y='total_bill', data=tips, ci='sd', palette='pastel')
plt.title('Average Total Bill per Day')
plt.show()
```

1.2.2 Boxplot and Violin Plot for Distribution Analysis

```
fig, axs = plt.subplots(1, 2, figsize=(12, 5))
sns.boxplot(x='day', y='total_bill', data=tips, ax=axs[0], palette='Blues')
axs[0].set_title('Boxplot: Total Bill per Day')
sns.violinplot(x='day', y='total_bill', data=tips, ax=axs[1], palette='Oranges')
axs[1].set_title('Violin Plot: Total Bill per Day')
plt.tight_layout()
plt.show()
```

1.3 Pairplot for Exploratory Data Analysis

Quickly explore relationships between numerical variables:

```
sns.pairplot(tips, hue='sex', palette='coolwarm')
plt.suptitle('Pairplot: Tips Dataset', y=1.02)
plt.show()
```

X 1.4 Heatmaps for Correlation Analysis

```
corr = tips.corr(numeric_only=True)
sns.heatmap(corr, annot=True, cmap='YlGnBu', linewidths=0.5)
plt.title('Correlation Heatmap')
plt.show()
```

3. Plotly for Interactive and Web-Based Visualizations

Plotly creates highly interactive and web-friendly visualizations with minimal code.

♦ 2.1 Installing Plotly

```
pip install plotly
```

2.2 Interactive Line and Bar Plots

Hover, zoom, and pan with ease in Plotly plots.

3.3 Choropleth Maps (Geospatial Visualizations)

11 2.4 3D Scatter Plots

(2) 3. Heatmaps, Correlation Matrices, and Geospatial Plots

Advanced analysis often requires specialized visualizations.

3.1 Complex Correlation Heatmaps

Using Seaborn with mask options for cleaner heatmaps:

```
import numpy as np

# Compute correlation
corr = tips.corr(numeric_only=True)

# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# Create heatmap
sns.heatmap(corr, mask=mask, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Upper Triangle Correlation Heatmap')
plt.show()
```

3.2 Geospatial Plots with Plotly and Geopandas

Geopandas allows manipulation of geospatial data, and Plotly enhances visualization.

4. 3D Plots, Network Graphs, and Complex Representations

(2) 4.1 Matplotlib 3D Plots

```
from mpl_toolkits.mplot3d import Axes3D

fig = plt.figure(figsize=(8,6))
ax = fig.add_subplot(111, projection='3d')

# Data
x = np.random.rand(100)
y = np.random.rand(100)
z = np.random.rand(100)

# 3D Scatter
ax.scatter(x, y, z, c=z, cmap='viridis')
ax.set_title('3D Scatter Plot')
plt.show()
```

4.2 Network Graphs with NetworkX

NetworkX enables graph-based data visualizations like social networks and relationships.

```
import networkx as nx

# Create graph
G = nx.erdos_renyi_graph(n=20, p=0.2)

# Plot graph
plt.figure(figsize=(10, 7))
nx.draw(G, with_labels=True, node_color='skyblue', node_size=700,
edge_color='gray')
plt.title('Random Network Graph')
plt.show()
```

4.3 Sankey Diagrams for Flow Visualizations (Plotly)

```
import plotly.graph_objects as go

fig = go.Figure(data=[go.Sankey(
    node=dict(label=["Source A", "Source B", "Target C", "Target D"]),
    link=dict(source=[0, 1, 0, 1],
        target=[2, 2, 3, 3],
        value=[8, 4, 2, 8]))])

fig.update_layout(title_text="Sankey Diagram Example", font_size=12)
fig.show()
```

✓ Chapter Summary

♠ Enhanced statistical visualizations with Seaborn ♠ Built interactive and web-based visualizations using
Plotly ♠ Explored heatmaps, correlation matrices, and geospatial plots ♠ Created 3D plots and network
graphs for complex data analysis

Chapter 12: Building Interactive Dashboards

In this chapter, we will explore how to create **interactive dashboards** that combine the power of **Pandas**, Matplotlib, and Plotly. We'll also integrate ipywidgets and Voila to add interactivity, and conclude with methods for creating reproducible reports using Jupyter and nbconvert.

By the end of this chapter, you'll be able to:

✓ Build dashboards that update dynamically based on user input ✓ Design visually engaging, interactive data applications <a> Convert Jupyter Notebooks into standalone dashboards



1. Integrating Pandas, Matplotlib, and Plotly for Dashboards

To build a powerful dashboard, we'll combine data processing (Pandas), static visualizations (Matplotlib), and interactive plots (Plotly).

1.1 Building a Simple Dashboard with Pandas & Matplotlib

Step 1: Prepare the Data

```
import pandas as pd
import matplotlib.pyplot as plt
# Sample dataset
df = pd.read_csv('https://raw.githubusercontent.com/mwaskom/seaborn-
data/master/iris.csv')
# Ouick data check
df.head()
```

Step 2: Create a Dynamic Plot

```
def plot species distribution(species):
   filtered_df = df[df['species'] == species]
   plt.figure(figsize=(8, 5))
   plt.hist(filtered_df['sepal_length'], bins=10, color='skyblue',
edgecolor='black')
   plt.title(f'Sepal Length Distribution for {species.capitalize()}')
   plt.xlabel('Sepal Length')
    plt.ylabel('Frequency')
    plt.show()
```

```
# Example call
plot_species_distribution('setosa')
```

4 1.2 Adding Interactivity with Plotly

Plotly enables zooming, panning, and tooltips without extra code.

2. Using ipywidgets and Voila for Interactivity

ipywidgets lets you create interactive controls, while Voila turns Jupyter notebooks into web apps.

2.1 Install ipywidgets and Voila

```
pip install ipywidgets voila
```

2.2 Creating Interactive Controls with ipywidgets

```
import ipywidgets as widgets
from IPython.display import display

# Dropdown widget
species_dropdown = widgets.Dropdown(
    options=df['species'].unique(),
    value='setosa',
    description='Species:',
)

# Function to update plot
def update_plot(species):
    plot_species_distribution(species)

# Link widget to function
widgets.interactive(update_plot, species=species_dropdown)
```

Run this cell and select different species from the dropdown to see live plot updates.

(#) 2.3 Deploying Dashboards with Voila

Voila removes code cells and only displays widgets and outputs.

Step 1: Save Your Notebook (e.g., iris_dashboard.ipynb)

Step 2: Run Voila

```
voila iris_dashboard.ipynb
```

This launches a browser-based dashboard that anyone can interact with.

3. Designing Reproducible Reports with Jupyter and nbconvert

To share your analysis as a report or presentation:

3.1 Install nbconvert

pip install nbconvert

3.2 Export as PDF or HTML

In Jupyter Notebook:

File → Download As → HTML/PDF

Or use terminal commands:

```
jupyter nbconvert --to html your_notebook.ipynb
jupyter nbconvert --to pdf your_notebook.ipynb
```

This creates polished, shareable reports directly from your analysis.

4. Building a Real-World Interactive Dashboard

4.1 Case Study: COVID-19 Data Dashboard

Step 1: Import Data

```
covid_url = 'https://covid.ourworldindata.org/data/owid-covid-data.csv'
covid_df = pd.read_csv(covid_url, parse_dates=['date'])

# Filter recent data
covid_df = covid_df[covid_df['date'] >= '2023-01-01']
```

Step 2: Create Interactive Country Selector

```
country_selector = widgets.Dropdown(
    options=covid_df['location'].unique(),
    value='United States',
    description='Country:',
)

def plot_covid_trend(country):
    data = covid_df[covid_df['location'] == country]
    fig = px.line(data, x='date', y='new_cases', title=f'COVID-19 Cases in
{country}')
    fig.show()

widgets.interactive(plot_covid_trend, country=country_selector)
```

Step 3: Deploy with Voila

```
voila covid_dashboard.ipynb
```

This creates a fully interactive COVID-19 dashboard that updates in real-time.

✓ Chapter Summary

- Duilt dynamic plots using Pandas, Matplotlib, and Plotly Created interactive widgets using ipywidgets
- Deployed dashboards as web apps using **Voila** Generated reproducible reports with **nbconvert**

oxdot Chapter 13: Optimizing Data Pipelines

In data analysis, working with large datasets often leads to bottlenecks in **performance** and **memory usage**. This chapter focuses on optimizing data pipelines for speed and efficiency. We'll cover profiling, vectorization, lazy evaluation, handling sparse data, and efficient disk I/O strategies to build highperformance data pipelines.

1. Profiling and Benchmarking Data Operations

Before optimizing code, it's essential to identify bottlenecks using **profiling** and **benchmarking** tools.

1.1 Profiling with cProfile and pstats

The cProfile module helps analyze code performance.

```
import cProfile
import pandas as pd
def process data():
    df = pd.DataFrame({'A': range(1000000), 'B': range(1000000)})
    df['C'] = df['A'] + df['B']
    return df
cProfile.run('process_data()')
```

For a more detailed report:

```
python -m cProfile -s time your_script.py
```

(i) 1.2 Benchmarking with %timeit in Jupyter

The **%timeit** magic command measures execution time with high precision.

```
import numpy as np
arr = np.random.rand(1 000 000)
%timeit np.sqrt(arr) # Fast, vectorized operation
```

1.3 Visualizing Bottlenecks with line_profiler

Install and use line_profiler to identify slow lines of code:

```
pip install line_profiler
```

```
%load_ext line_profiler

def slow_function():
    total = 0
    for i in range(1000000):
        total += i
    return total

%lprun -f slow_function slow_function()
```

2. Performance Tuning with Vectorization and Lazy Evaluation

2.1 Vectorization: Replacing Loops with Numpy Operations

Vectorized operations in **Numpy** and **Pandas** avoid Python-level loops, dramatically improving performance.

X Inefficient (Using Loops):

```
arr = np.random.rand(1_000_000)
result = []

for val in arr:
    result.append(val ** 2)
```

Efficient (Vectorized):

```
result = arr ** 2 # ~100x faster
```

四 2.2 Lazy Evaluation with Dask

Dask enables lazy evaluation for larger-than-memory datasets.

```
pip install dask
```

```
import dask.dataframe as dd

# Load a large CSV with Dask
df = dd.read_csv('large_dataset.csv')

# Lazy computation (not yet executed)
result = df.groupby('category')['value'].mean()

# Trigger computation
result.compute()
```

Key Benefits: Handles out-of-core datasets Parallelizes operations across cores

2.3 Applying Vectorization in Pandas

Use built-in Pandas functions instead of apply or loops:

X Inefficient:

```
df['new_col'] = df['value'].apply(lambda x: x * 2)
```

Efficient:

```
df['new_col'] = df['value'] * 2
```

3. Handling Sparse Data and Optimizing Data Types

Reducing memory usage by optimizing column types.

```
df = pd.read_csv('large_dataset.csv')

# Check memory usage
print(df.info(memory_usage='deep'))

# Convert float64 to float32 and int64 to int32
df['int_column'] = df['int_column'].astype('int32')
df['float_column'] = df['float_column'].astype('float32')
```

Tip: Use category type for low-cardinality text columns.

```
df['category_column'] = df['category_column'].astype('category')
```

3.2 Using Sparse Matrices for High-Dimensional Data

Scipy's sparse matrices help when most data points are zeros.

```
from scipy import sparse
import numpy as np

dense_matrix = np.random.randint(0, 2, size=(1000, 1000))
sparse_matrix = sparse.csr_matrix(dense_matrix)

print(f"Memory usage of dense matrix: {dense_matrix.nbytes / 1e6} MB")
print(f"Memory usage of sparse matrix: {sparse_matrix.data.nbytes / 1e6} MB")
```

Sparse matrices significantly reduce memory usage in cases of sparse data.

4. Memory Mapping and Efficient Disk I/O

3 4.1 Memory-Mapped Files with Numpy

Memory mapping enables access to data on disk as if it were in memory without loading it all at once.

```
import numpy as np

# Create a memory-mapped file
data = np.memmap('data.dat', dtype='float32', mode='w+', shape=(10000, 10000))
data[:] = np.random.rand(10000, 10000)
del data # Flush to disk

# Access a subset without loading the entire file
mapped_data = np.memmap('data.dat', dtype='float32', mode='r', shape=(10000, 10000))
print(mapped_data[5000, 5000]) # Access specific data point
```

4.2 Efficient File Formats: Parquet and Feather

For large-scale data storage and retrieval, use columnar storage formats like Parquet and Feather.

```
pip install pyarrow fastparquet
```

```
# Save DataFrame as Parquet
df.to_parquet('data.parquet', engine='pyarrow')

# Read Parquet (faster than CSV)
df = pd.read_parquet('data.parquet', engine='pyarrow')
```

Benefits: ✓ Faster read/write times ✓ Built-in compression ✓ Optimized for column-wise operations

☐ 4.3 Chunking Large Files for Out-of-Core Processing

When dealing with massive CSV files:

```
chunksize = 10 ** 6  # Read 1 million rows at a time

for chunk in pd.read_csv('large_file.csv', chunksize=chunksize):
    # Perform aggregation on each chunk
    result = chunk.groupby('category')['value'].sum()
    print(result.head())
```

Tip: Combine with Dask for parallel processing.

Chapter Summary

♦ Identified bottlenecks using **profiling** and **benchmarking** ♦ Improved performance with **vectorization** and **lazy evaluation** ♦ Reduced memory usage by **optimizing data types** and using **sparse matrices** ♦ Enhanced I/O efficiency with **memory mapping** and **columnar file formats**

Chapter 14: Memory Management and Scalability

As datasets grow in size and complexity, efficient **memory management** and **scalability** become critical in data analysis. This chapter focuses on techniques for reducing memory footprints, working with **large-scale datasets** using tools like **Dask** and **PySpark**, implementing **parallel and distributed data analysis**, and handling **real-time analytics** using **caching** and **streaming data**.

By the end of this chapter, you will be able to:

- Optimize memory usage in Pandas and Numpy Work with massive datasets using Dask and PySpark
- Apply parallel and distributed processing techniques <a>Implement real-time analytics with caching and streaming

1. Reducing Memory Footprint in Pandas and Numpy

Efficient memory management is key when working with large datasets in **Pandas** and **Numpy**.

1.1 Memory Optimization in Pandas

Step 1: Analyze Memory Usage

```
import pandas as pd

# Load dataset

df = pd.read_csv('large_dataset.csv')

# Check memory usage

df.info(memory_usage='deep')
```

♦ Step 2: Optimize Data Types

- Convert floats to float32 if precision allows.
- Convert integers to smaller types like int16 or int32.
- Use categorical types for repetitive text data.

```
# Optimize numeric columns
df['int_column'] = pd.to_numeric(df['int_column'], downcast='integer')
df['float_column'] = pd.to_numeric(df['float_column'], downcast='float')

# Optimize string columns
df['category_column'] = df['category_column'].astype('category')
```

```
# Check memory again
df.info(memory_usage='deep')
```

Tip: These optimizations can reduce memory usage by 50-90%.

1.2 Efficient Memory Usage in Numpy

Numpy arrays can be optimized by choosing the right data types:

```
import numpy as np

# Default: float64
arr = np.random.rand(1_000_000)

# Reduce memory with float32
arr_32 = arr.astype('float32')

print(f"Original size: {arr.nbytes / 1e6:.2f} MB")
print(f"Optimized size: {arr_32.nbytes / 1e6:.2f} MB")
```

Result: Float32 reduces memory by **50%** with minimal precision loss.

2. Working with Large-Scale Datasets Using Dask and PySpark

When data exceeds your machine's memory, use distributed frameworks like Dask and PySpark.

1 2.1 Scaling Pandas with Dask

Dask extends Pandas for parallel and out-of-core computations.

Install Dask:

```
pip install dask[complete]
```

Example: Using Dask DataFrames

```
import dask.dataframe as dd

# Load large CSV with Dask
df = dd.read_csv('large_dataset.csv')

# Lazy operations
result = df.groupby('category')['value'].mean()
```

```
# Execute computation
result.compute()
```

© Key Benefits: ✓ Handles datasets larger than RAM ✓ Parallelizes across multiple CPU cores ✓ Syntax similar to Pandas

(2.2 Distributed Computing with PySpark

PySpark is a Python API for Apache Spark, designed for distributed data analysis across clusters.

Install PySpark:

```
pip install pyspark
```

Example: PySpark DataFrame

```
from pyspark.sql import SparkSession

# Initialize Spark
spark = SparkSession.builder.appName("BigDataApp").getOrCreate()

# Load data
df = spark.read.csv('large_dataset.csv', header=True, inferSchema=True)

# Group and aggregate
df.groupBy('category').avg('value').show()
```

Why PySpark? ✓ Handles terabyte-scale datasets ✓ Built for distributed environments ✓ Supports SQL-like queries, machine learning, and graph processing

3. Parallel and Distributed Data Analysis Techniques

3.1 Parallel Processing in Pandas Using joblib

For parallelizing custom functions in Pandas:

```
pip install joblib
```

```
from joblib import Parallel, delayed import pandas as pd
```

```
df = pd.DataFrame({'A': range(1000000)})

# Parallel apply
def process_row(x):
    return x ** 2

results = Parallel(n_jobs=-1)(delayed(process_row)(x) for x in df['A'])
```

▽ Tip: n_jobs=-1 uses all available CPU cores.

III 3.2 Distributed Data Processing with Dask and PySpark

Framework	Best For	Strength
Dask	Medium datasets, Pythonic	Easy Pandas migration
PySpark	Big data, distributed jobs	Cluster-level parallelization

4. Caching, Streaming Data, and Real-Time Analytics

For real-time data analysis and low-latency applications, leverage caching and streaming frameworks.

4.1 Caching Data for Speed

Pandas and Dask can cache intermediate results to avoid redundant computations.

In Dask:

```
import dask.cache

# Enable cache
cache = dask.cache.Cache(1e9) # 1 GB cache
cache.register()

# Cached computation
df = dd.read_csv('large_dataset.csv').persist()
```

4.2 Streaming Data with PySpark Structured Streaming

Handle real-time data streams using PySpark:

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("StreamingApp").getOrCreate()
```

Q Use Cases: ✓ IoT device streams ✓ Real-time dashboards ✓ Log file monitoring

4.3 Real-Time Analytics with Kafka and Spark

For complex real-time systems, combine **Kafka** (for data streams) with **Spark** (for processing).

Pipeline Example:

- 1. Kafka streams data →
- 2. Spark processes data →
- 3. Dashboards update in real-time

Chapter Summary

♠ Reduced memory usage with efficient data types and sparse matrices ♦ Scaled data pipelines using **Dask** and **PySpark** ♦ Applied parallel and distributed data analysis techniques ♦ Implemented real-time analytics using **caching** and **streaming**

Chapter 15: Advanced Data Wrangling and Transformation

In large-scale data analysis, mastering **data wrangling** and **transformation** is essential to extract meaningful insights. This chapter focuses on advanced techniques for reshaping data, performing complex aggregations, working with time series, and handling high-dimensional data.

By the end of this chapter, you will be able to:

Reshape complex datasets using **melt**, **pivot**, **stack**, and **unstack** Apply advanced **GroupBy** operations, **rolling windows**, and **window functions** Perform **time series analysis** and detect anomalies Manage **high-dimensional** and **multivariate** data

1. Complex Reshaping Techniques (Melt, Pivot, Stack, Unstack)

Reshaping data is key to aligning it with specific analysis needs. Pandas provides versatile tools like **melt**, **pivot**, **stack**, and **unstack**.

1.1 Melting DataFrames

The melt function transforms wide-format data into long-format, making it suitable for analysis.

Example:

```
import pandas as pd

df = pd.DataFrame({
    'Product': ['A', 'B', 'C'],
    'Jan_Sales': [200, 150, 300],
    'Feb_Sales': [220, 180, 320]
})

# Melt the DataFrame
melted_df = pd.melt(df, id_vars=['Product'], var_name='Month', value_name='Sales')

print(melted_df)
```

Output:

```
Product Month Sales

0 A Jan_Sales 200

1 B Jan_Sales 150

2 C Jan_Sales 300

3 A Feb_Sales 220
```

```
4 B Feb_Sales 180
5 C Feb_Sales 320
```

1.2 Pivoting DataFrames

The pivot method reverses the effect of **melt**, turning long-format data back into wide format.

```
pivot_df = melted_df.pivot(index='Product', columns='Month', values='Sales')
print(pivot_df)
```

Output:

```
Month Jan_Sales Feb_Sales
Product
A 200 220
B 150 180
C 300 320
```

1.3 Stacking and Unstacking

- stack() converts columns into rows (creates a hierarchical index).
- unstack() moves rows back to columns.

Example:

```
df = pd.DataFrame({
    'City': ['NY', 'LA', 'SF'],
    '2023_Sales': [500, 600, 550],
    '2024_Sales': [520, 610, 580]
}).set_index('City')

# Stack: Convert columns into rows
stacked = df.stack()
print(stacked)

# Unstack: Reverse operation
unstacked = stacked.unstack()
print(unstacked)
```

2. Advanced GroupBy, Rolling Windows, and Window Functions

2.1 Complex GroupBy Operations

Use **groupby** with custom functions and multi-level aggregations.

Example: Multiple Aggregations

Output:

```
Salary Bonus
mean max sum

Department
Finance 81000.0 82000 16500
HR 52500.0 55000 11000
IT 71000.0 72000 14500
```

2.2 Rolling Windows and Expanding Windows

Rolling windows and **expanding windows** help analyze time-series trends.

Example: Rolling Mean

```
dates = pd.date_range('2024-01-01', periods=10)
df = pd.DataFrame({'Date': dates, 'Sales': [5, 6, 7, 8, 7, 6, 5, 6, 7, 8]})
df.set_index('Date', inplace=True)

# Apply a rolling window
df['Rolling_Avg'] = df['Sales'].rolling(window=3).mean()
print(df)
```

Expanding Window:

```
df['Expanding_Mean'] = df['Sales'].expanding().mean()
```

2.3 Window Functions with groupby

Apply window functions like rank, lead/lag, and cumsum.

Example: Cumulative Sum

```
df['Cumulative_Sum'] = df['Sales'].cumsum()
print(df)
```

3. Time Series Analysis and Anomaly Detection

31 3.1 Time Series Manipulation in Pandas

Use resample, shift, and rolling for time-based aggregations.

Example: Monthly Resampling

```
date_rng = pd.date_range(start='2024-01-01', end='2024-12-31', freq='D')
df = pd.DataFrame(date_rng, columns=['date'])
df['sales'] = np.random.randint(100, 500, size=(len(date_rng)))
df.set_index('date', inplace=True)

# Resample to monthly sales
monthly_sales = df.resample('M').sum()
print(monthly_sales.head())
```


Detect outliers using the **Z-Score** method.

```
from scipy import stats

df['z_score'] = stats.zscore(df['sales'])

# Flag anomalies

df['anomaly'] = df['z_score'].apply(lambda x: 'Anomaly' if abs(x) > 2 else
'Normal')
```

✓ 3.3 Seasonal Decomposition with statsmodels

```
pip install statsmodels
```

```
from statsmodels.tsa.seasonal import seasonal_decompose
result = seasonal_decompose(df['sales'], model='additive', period=30)
result.plot();
```

This decomposes time series into **trend**, **seasonality**, and **residuals**.

4. Handling High-Dimensional and Multivariate Data

4.1 Reducing Dimensionality with PCA

For high-dimensional data, Principal Component Analysis (PCA) simplifies datasets while retaining most variance.

```
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import numpy as np
# Simulated data
data = np.random.rand(100, 10)
# Standardize
scaler = StandardScaler()
data_scaled = scaler.fit_transform(data)
# Apply PCA
pca = PCA(n_components=2)
principal_components = pca.fit_transform(data_scaled)
print(principal_components[:5]) # First 5 data points
```

\$4.2 Handling Multivariate Time Series

Use multi-indexing and panel data analysis for time series with multiple variables.

Example:

```
arrays = [
    ['Region_A', 'Region_B', 'Region_B'],
    pd.date_range('2024-01-01', periods=2).tolist() * 2
]

index = pd.MultiIndex.from_tuples(list(zip(*arrays)), names=['Region', 'Date'])

df = pd.DataFrame({'Sales': [100, 120, 150, 160]}, index=index)

# Access specific region data
print(df.loc['Region_A'])
```

✓ Chapter Summary

♠ Reshaped complex datasets using melt, pivot, and stack/unstack ♠ Applied advanced GroupBy techniques, rolling windows, and window functions ♠ Conducted time series analysis and detected anomalies ♠ Managed high-dimensional data using PCA and multi-indexing

Chapter 16: Dimensionality Reduction and Complex Data Analysis

High-dimensional datasets can be overwhelming and challenging to interpret. Dimensionality reduction simplifies such data without significant loss of information, making it easier to visualize, analyze, and use in predictive models. This chapter covers essential techniques like **PCA**, **SVD**, **t-SNE**, and **UMAP**, along with methods for **feature engineering** and handling **multicollinearity**.

By the end of this chapter, you will be able to:

Apply **PCA** and **SVD** for dimensionality reduction Use **t-SNE** and **UMAP** for data visualization Perform effective **feature engineering** for predictive analytics Analyze **correlations** and manage **multicollinearity**

1. Principal Component Analysis (PCA) and Singular Value Decomposition (SVD)

III 1.1 Principal Component Analysis (PCA)

PCA is a linear dimensionality reduction technique that transforms data into new axes (principal components) that capture the most variance.

Example: Applying PCA

```
import pandas as pd
import numpy as np
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
# Sample dataset
np.random.seed(42)
data = np.random.rand(100, 5) * 100
df = pd.DataFrame(data, columns=['Feature1', 'Feature2', 'Feature3', 'Feature4',
'Feature5'])
# Standardizing the data
scaler = StandardScaler()
scaled_data = scaler.fit_transform(df)
# Apply PCA
pca = PCA(n_components=2)
principal components = pca.fit transform(scaled data)
# Create DataFrame for visualization
pca df = pd.DataFrame(data=principal components, columns=['PC1', 'PC2'])
```

```
# Plot the results
plt.figure(figsize=(8, 6))
plt.scatter(pca_df['PC1'], pca_df['PC2'], alpha=0.7)
plt.title('PCA - First Two Principal Components')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```

∜ Key Points:

- PCA reduces noise and redundancy.
- The explained variance ratio helps determine how much information is retained.

```
print("Explained Variance Ratio:", pca.explained_variance_ratio_)
```

1.2 Singular Value Decomposition (SVD)

SVD is a matrix factorization method that breaks down a matrix into singular vectors and values. It's commonly used in PCA and recommendation systems.

Example: Using SVD

```
from numpy.linalg import svd

# Random matrix
A = np.random.rand(5, 5)

# Perform SVD
U, Sigma, Vt = svd(A)

print("U Matrix:\n", U)
print("\nSigma Values:\n", Sigma)
print("\nV Transposed:\n", Vt)
```

₩ When to Use PCA vs. SVD:

- **PCA** is optimized for variance preservation.
- **SVD** is versatile for dimensionality reduction, text mining (LSA), and matrix compression.

2. t-SNE, UMAP, and Other Visualization Techniques

For non-linear data structures, **t-SNE** and **UMAP** provide better clustering and visualization.

2.1 t-SNE (t-Distributed Stochastic Neighbor Embedding)

t-SNE is a non-linear technique for high-dimensional data visualization in 2D or 3D.

Example: Visualizing Clusters with t-SNE

```
from sklearn.manifold import TSNE

# Apply t-SNE
tsne = TSNE(n_components=2, perplexity=30, random_state=42)
tsne_results = tsne.fit_transform(scaled_data)

# Plot
plt.figure(figsize=(8, 6))
plt.scatter(tsne_results[:, 0], tsne_results[:, 1], c='blue', alpha=0.7)
plt.title('t-SNE Visualization')
plt.xlabel('Dimension 1')
plt.ylabel('Dimension 2')
plt.show()
```

♣ Pros & Cons:

- Great for visualizing clusters.
- X Computationally expensive on large datasets.

② 2.2 UMAP (Uniform Manifold Approximation and Projection)

UMAP preserves more global structure than t-SNE and is faster for large datasets.

Example: Using UMAP

```
import umap

# Apply UMAP

umap_reducer = umap.UMAP(n_neighbors=15, min_dist=0.1, random_state=42)

umap_results = umap_reducer.fit_transform(scaled_data)

# Plot

plt.figure(figsize=(8, 6))

plt.scatter(umap_results[:, 0], umap_results[:, 1], c='green', alpha=0.7)

plt.title('UMAP Visualization')

plt.xlabel('UMAP1')

plt.ylabel('UMAP2')

plt.show()
```


- t-SNE focuses on local relationships.
- UMAP preserves both local and global data structures.

% 3. Feature Engineering for Predictive Analytics

Feature engineering transforms raw data into meaningful features that improve model performance.

3.1 Creating New Features

- Polynomial Features: Capture non-linear relationships.
- **Aggregations**: Mean, sum, counts for groupings.
- Domain Knowledge: Derive features based on specific problem contexts.

Example: Polynomial Features

```
from sklearn.preprocessing import PolynomialFeatures

# Create interaction terms and polynomial features
poly = PolynomialFeatures(degree=2, include_bias=False)
poly_features = poly.fit_transform(df[['Feature1', 'Feature2']])

print("Polynomial Features Shape:", poly_features.shape)
```

3.2 Scaling and Normalization

- StandardScaler for Gaussian-like data.
- MinMaxScaler for bounded ranges.
- RobustScaler for outlier-heavy data.

```
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
normalized_data = scaler.fit_transform(df)
```

3.3 Encoding Categorical Variables

- One-Hot Encoding for non-ordinal categories.
- Ordinal Encoding for ordered variables.

```
df = pd.DataFrame({'Category': ['A', 'B', 'C', 'A', 'B']})
encoded_df = pd.get_dummies(df, columns=['Category'])
print(encoded_df)
```

III 4. Correlation Analysis and Multicollinearity

Understanding relationships between variables is crucial in data analysis and modeling.

4.1 Correlation Analysis

Use correlation matrices and heatmaps to visualize variable relationships.

Example: Correlation Heatmap

```
import seaborn as sns

# Correlation matrix
corr_matrix = df.corr()

# Plot heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Feature Correlation Heatmap')
plt.show()
```


Multicollinearity (high correlation between independent variables) can distort regression models.

Detection Methods:

• Variance Inflation Factor (VIF): Detects multicollinearity.

```
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Compute VIF for each feature
vif_data = pd.DataFrame()
vif_data["feature"] = df.columns
vif_data["VIF"] = [variance_inflation_factor(scaled_data, i) for i in
range(scaled_data.shape[1])]
print(vif_data)
```

✓ Solution Strategies:

- Remove highly correlated features.
- Apply dimensionality reduction (PCA).
- Use regularization techniques (Ridge, Lasso).

✓ Chapter Summary

♦ Applied **PCA** and **SVD** for linear dimensionality reduction ♦ Used **t-SNE** and **UMAP** for non-linear data visualization ♦ Performed **feature engineering** for predictive modeling ♦ Analyzed **correlations** and addressed **multicollinearity**

☐ Chapter 17: Exploratory Data Analysis (EDA) in Practice

Exploratory Data Analysis (EDA) is a critical step in the data analysis pipeline. It helps analysts uncover patterns, detect anomalies, test hypotheses, and check assumptions using statistical summaries and visualizations. In this chapter, we will build effective EDA workflows, automate the process using custom Python scripts, and apply EDA techniques to a real-world case study.

By the end of this chapter, you will be able to:

Design efficient **EDA** workflows Identify patterns, trends, and anomalies Automate **EDA** using Pythonic techniques Apply EDA in a real-world case study

① 1. Designing Effective EDA Workflows

A structured EDA process helps transform raw data into actionable insights. Here's a typical EDA workflow:

- 1. Data Collection & Loading
- 2. Data Profiling & Quality Checks
- 3. Data Cleaning & Preprocessing
- 4. Exploratory Visualizations
- 5. Feature Engineering & Transformation
- 6. Pattern & Anomaly Detection
- 7. Insights Summary

4 1.1 Essential Python Libraries for EDA

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from pandas_profiling import ProfileReport
```

1.2 Quick Data Overview

```
# Sample E-commerce Dataset
df = pd.read_csv('ecommerce_data.csv')

# Quick data inspection
print(df.head())
print(df.info())
print(df.describe())
```

- df.info() → Data types, null values, memory usage
- df.describe() → Statistical summary of numerical features
- **▼ Tip:** Use pandas_profiling for quick automated profiling.

```
profile = ProfileReport(df, title="E-Commerce Data Report", explorative=True)
profile.to_file("ecommerce_eda_report.html")
```

11 2. Identifying Patterns, Trends, and Anomalies

Exploratory Visualizations help reveal hidden insights.

2.1 Univariate Analysis

Examine individual features.

```
# Distribution of Purchase Amount
sns.histplot(df['PurchaseAmount'], bins=30, kde=True)
plt.title('Distribution of Purchase Amount')
plt.show()
```

- **Histograms** → Distribution of numerical data
- **Bar plots** → Categorical data frequency

1 2.2 Bivariate Analysis

Understand relationships between variables.

```
# Purchase Amount vs. Customer Age
sns.scatterplot(x='Age', y='PurchaseAmount', data=df)
plt.title('Age vs. Purchase Amount')
plt.show()

# Correlation Heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title('Feature Correlation Matrix')
plt.show()
```

- **Scatterplots** → Continuous relationships
- **Heatmaps** → Correlations and multicollinearity

⚠ 2.3 Anomaly Detection

Outliers can distort analysis and predictions.

```
# Boxplot to detect outliers
sns.boxplot(x=df['PurchaseAmount'])
plt.title('Outlier Detection in Purchase Amount')
plt.show()
```

Advanced Techniques for Anomaly Detection:

- Z-Score Method
- IQR Method
- Isolation Forest & DBSCAN (for complex data)

```
# Apply Isolation Forest
iso_forest = IsolationForest(contamination=0.05)
df['anomaly'] = iso_forest.fit_predict(df[['PurchaseAmount', 'Age']])
sns.scatterplot(x='Age', y='PurchaseAmount', hue='anomaly', data=df)
plt.title('Anomaly Detection with Isolation Forest')
plt.show()
```

3. Automating EDA with Custom Scripts

For repetitive EDA tasks, Pythonic automation improves efficiency.

3.1 Reusable EDA Function

```
def automated_eda(dataframe):
    print("Shape of Dataset:", dataframe.shape)
    print("\nMissing Values:\n", dataframe.isnull().sum())
    print("\nData Types:\n", dataframe.dtypes)
    print("\nStatistical Summary:\n", dataframe.describe())

# Correlation Heatmap
    plt.figure(figsize=(10, 6))
    sns.heatmap(dataframe.corr(), annot=True, cmap='Blues')
    plt.title('Correlation Matrix')
    plt.show()

# Run EDA
automated_eda(df)
```

% 3.2 Generating Automated Reports

```
# Full profiling report
profile = ProfileReport(df, title="Automated EDA Report", explorative=True)
profile.to_file("automated_eda_report.html")
```

☐ 4. Case Study: Analyzing E-Commerce Customer Behavior

Dataset: Online retail transactions including demographics, purchase behavior, and browsing patterns.

Objective: Identify customer segments and factors affecting purchase amounts.

3 4.1 Data Overview

```
# Load dataset
df = pd.read_csv('ecommerce_data.csv')
# Check for null values and data types
print(df.info())
```

Ⅲ 4.2 Key EDA Insights

- 1. Who are the high-value customers?
- 2. Which products are most popular?
- 3. Are there seasonal trends in purchase patterns?

4.3 Visualizing Customer Behavior

```
# Customer Segmentation by Age and Purchase Amount
sns.scatterplot(x='Age', y='PurchaseAmount', hue='CustomerSegment', data=df)
plt.title('Customer Segments: Age vs. Purchase Amount')
plt.show()

# Monthly Sales Trend
df['OrderDate'] = pd.to_datetime(df['OrderDate'])
df['Month'] = df['OrderDate'].dt.to_period('M')
monthly_sales = df.groupby('Month')['PurchaseAmount'].sum().reset_index()

plt.figure(figsize=(12, 6))
sns.lineplot(x='Month', y='PurchaseAmount', data=monthly_sales)
plt.title('Monthly Sales Trend')
plt.show()
```

4.4 RFM (Recency, Frequency, Monetary) Analysis

RFM Analysis segments customers based on:

- Recency (How recently a customer made a purchase)
- **Frequency** (How often they purchase)
- Monetary (How much they spend)

Visualization:

```
# RFM Distribution
fig, ax = plt.subplots(1, 3, figsize=(18, 5))
sns.histplot(rfm['Recency'], bins=30, ax=ax[0])
sns.histplot(rfm['Frequency'], bins=30, ax=ax[1])
sns.histplot(rfm['Monetary'], bins=30, ax=ax[2])
ax[0].set_title('Recency Distribution')
ax[1].set_title('Frequency Distribution')
ax[2].set_title('Monetary Distribution')
plt.show()
```

Chapter Summary

◆ Designed structured EDA workflows ◆ Explored patterns, trends, and detected anomalies ◆ Automated EDA tasks using Pythonic techniques ◆ Applied EDA in a real-world e-commerce case study

Chapter 18: Domain-Specific Data Analysis

In data analysis, domain knowledge shapes how we process, analyze, and interpret data. This chapter explores domain-specific data analysis techniques across finance, healthcare, social media, and scientific computing. We'll apply **Python libraries** like pandas, numpy, and visualization tools while integrating **domain-specific methodologies**.

By the end of this chapter, you'll be able to:

Apply time series analysis for **financial data** Use statistical methods for **healthcare analytics** Perform text-based **sentiment analysis** Conduct **scientific simulations** using Monte Carlo methods

(§) 1. Financial Data Analysis

Financial data often involves **time series**, where the order of data points is critical. Analysts focus on trends, seasonality, volatility, and risk.

III 1.1 Time Series Analysis

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import yfinance as yf

# Load historical stock data
stock_data = yf.download('AAPL', start='2022-01-01', end='2023-12-31')
stock_data['Close'].plot(title='AAPL Stock Prices', figsize=(10, 5))
plt.show()
```

1.2 Moving Averages & Volatility

```
# 20-day Moving Average
stock_data['20_MA'] = stock_data['Close'].rolling(window=20).mean()

# Bollinger Bands
stock_data['Upper'] = stock_data['20_MA'] + 2 *
stock_data['Close'].rolling(window=20).std()
stock_data['Lower'] = stock_data['20_MA'] - 2 *
stock_data['Close'].rolling(window=20).std()

# Visualization
plt.figure(figsize=(12, 6))
plt.plot(stock_data['Close'], label='Closing Price')
plt.plot(stock_data['20_MA'], label='20-Day MA', color='orange')
plt.fill_between(stock_data.index, stock_data['Upper'], stock_data['Lower'],
color='gray', alpha=0.3)
```

```
plt.legend()
plt.title('AAPL Stock with Bollinger Bands')
plt.show()
```



```
from statsmodels.tsa.arima.model import ARIMA

# ARIMA Model
model = ARIMA(stock_data['Close'], order=(5, 1, 0))
model_fit = model.fit()
forecast = model_fit.forecast(steps=30)

# Plot Forecast
plt.figure(figsize=(10, 5))
plt.plot(stock_data['Close'], label='Actual')
plt.plot(pd.date_range(start=stock_data.index[-1], periods=30, freq='D'),
forecast, label='Forecast', color='red')
plt.legend()
plt.title('AAPL Stock Price Forecast')
plt.show()
```

4 1.4 Risk Modeling: Value at Risk (VaR)

```
# Historical VaR at 95% confidence
daily_returns = stock_data['Close'].pct_change().dropna()
var_95 = np.percentile(daily_returns, 5)
print(f"95% Value at Risk (VaR): {var_95 * 100:.2f}%")
```

2. Healthcare Analytics

Healthcare data requires handling sensitive, high-dimensional data with a focus on statistical accuracy.

11 2.1 Patient Data Analysis

```
# Sample healthcare data
df_health = pd.read_csv('patient_data.csv')

# Average patient age by disease type
df_health.groupby('Disease')['Age'].mean().plot(kind='bar', title='Average Age by
Disease')
plt.show()
```

∠ 2.2 Survival Analysis (Kaplan-Meier Curve)

```
from lifelines import KaplanMeierFitter

# Sample survival data

df_survival = pd.DataFrame({
    'duration': [5, 6, 6, 2, 4, 4, 3],
    'event_observed': [1, 0, 0, 1, 1, 1, 0]
})

kmf = KaplanMeierFitter()
kmf.fit(df_survival['duration'], event_observed=df_survival['event_observed'])

kmf.plot()
plt.title('Kaplan-Meier Survival Curve')
plt.show()
```

2.3 Predictive Modeling for Disease Diagnosis

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Prepare data
X = df_health.drop('Diagnosis', axis=1)
y = df_health['Diagnosis']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Random Forest Classifier
model = RandomForestClassifier()
model.fit(X_train, y_train)
predictions = model.predict(X_test)

# Accuracy
print("Accuracy:", accuracy_score(y_test, predictions))
```

3. Social Media & Text Data Analytics

Social media data is mostly **unstructured text**. Analyzing user sentiment and trends provides valuable insights.

III 3.1 Text Preprocessing

```
import pandas as pd
import re
from nltk.corpus import stopwords
```

(2) 3.2 Sentiment Analysis Using NLTK

```
from nltk.sentiment import SentimentIntensityAnalyzer

sia = SentimentIntensityAnalyzer()
tweets['sentiment_score'] = tweets['cleaned'].apply(lambda x:
sia.polarity_scores(x)['compound'])

# Classify sentiments
tweets['sentiment'] = tweets['sentiment_score'].apply(lambda x: 'Positive' if x >
0.05 else ('Negative' if x < -0.05 else 'Neutral'))
print(tweets[['text', 'sentiment']])</pre>
```

☑ 3.3 Word Cloud for Trend Analysis

```
from wordcloud import WordCloud

# Combine all tweets
all_text = ' '.join(tweets['cleaned'])

# Generate word cloud
wordcloud = WordCloud(width=800, height=400,
background_color='white').generate(all_text)

plt.figure(figsize=(10, 5))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title('Social Media Word Cloud')
plt.show()
```

Monte Carlo simulations use **random sampling** to solve complex problems that are deterministic in nature but difficult to solve directly.

4.1 Estimating Pi with Monte Carlo

```
import numpy as np
import matplotlib.pyplot as plt
# Simulation Parameters
n points = 10000
circle_points = 0
# Generate random points
x = np.random.rand(n_points)
y = np.random.rand(n_points)
# Calculate distance from origin
distance = np.sqrt(x**2 + y**2)
circle_points = np.sum(distance <= 1)</pre>
# Estimate Pi
pi_estimate = (circle_points / n_points) * 4
print(f"Estimated Pi: {pi_estimate}")
# Visualization
plt.figure(figsize=(6, 6))
plt.scatter(x[distance <= 1], y[distance <= 1], color='blue', s=1, label='Inside</pre>
Circle')
plt.scatter(x[distance > 1], y[distance > 1], color='red', s=1, label='Outside
Circle')
plt.legend()
plt.title('Monte Carlo Estimation of Pi')
plt.show()
```

4.2 Risk Analysis Using Monte Carlo

```
# Simulate stock returns
simulations = []
for _ in range(1000):
    returns = np.random.normal(0.001, 0.02, 252)
    cumulative_return = np.prod(1 + returns) - 1
    simulations.append(cumulative_return)

# Plot results
plt.hist(simulations, bins=50, edgecolor='black')
plt.title('Monte Carlo Simulation of Annual Returns')
plt.xlabel('Return')
plt.ylabel('Frequency')
plt.show()
```

✓ Chapter Summary

♦ Explored time series analysis and risk modeling in finance
 ♦ Applied statistical methods and
 predictive modeling in healthcare
 ♦ Performed text analysis and sentiment classification for social media
 ♦ Used Monte Carlo simulations for complex scientific problems

☐ Chapter 19: Building Scalable Data Analysis Pipelines

Modern data projects require scalability, reproducibility, and maintainability. This chapter focuses on building scalable data analysis pipelines using best practices for structuring complex projects, implementing data versioning, automating ETL workflows, and deploying end-to-end pipelines.

By the end of this chapter, you'll be able to:

Structure complex data projects for scalability and reusability Apply data versioning with DVC and Git Build automated ETL (Extract, Transform, Load) pipelines Implement a real-world churn prediction pipeline

1. Structuring Complex Data Projects for Reusability

A well-structured project ensures maintainability, reproducibility, and collaboration.

1.1 Recommended Project Structure

```
churn-prediction/
                     # Raw and processed data
  – data/
   - raw/
    ___ processed/
  - notebooks/
              # Jupyter notebooks for exploration
  - src/
                       # Source code for data pipelines
   ├─ data_ingestion.py
    — data preprocessing.py
     feature_engineering.py
    └─ model_training.py
  - models/
                      # Trained models
  - reports/
                      # Generated reports and visualizations
  - requirements.txt  # Dependencies
  - dvc.yaml
                      # Data pipeline configuration (DVC)
  - README.md
```

% 1.2 Best Practices

- Modular Code: Split code into reusable modules (ingestion, preprocessing, modeling).
- **Environment Management:** Use requirements.txt or conda for dependency control.
- Reproducibility: Leverage version control for code (Git) and data (DVC).

• Automated Workflows: Use tools like Airflow or Prefect for orchestrating tasks.

11 2. Data Versioning and Reproducibility with DVC and Git

Traditional Git struggles with large datasets. **DVC (Data Version Control)** tracks datasets and models, ensuring reproducibility.

♣ 2.1 Setting Up DVC

1. Install DVC:

```
pip install dvc
```

2. Initialize DVC in the Project:

```
git init
dvc init
```

3. Track Data Files:

```
dvc add data/raw/customer_data.csv
git add data/.gitignore data/raw/customer_data.csv.dvc
git commit -m "Track raw customer data with DVC"
```

4. Remote Storage (optional):

```
dvc remote add -d myremote s3://mybucket/mydata
dvc push
```

2.2 Managing Data Versions

- Modify Data: Update your dataset (data/raw/customer_data.csv).
- Track Changes:

```
dvc add data/raw/customer_data.csv
git commit -am "Updated customer data"
dvc push
```

• Switch Between Versions:

```
git checkout previous_commit_hash
dvc checkout
```

Tip: Combine Git (code) + DVC (data) for complete project reproducibility.

3. Creating Automated ETL Pipelines

An ETL Pipeline automates the process of Extracting, Transforming, and Loading data.

3.1 Designing the ETL Flow

- 1. Extract: Pull data from CSV, APIs, or databases.
- 2. Transform: Clean, preprocess, and engineer features.
- 3. **Load:** Save processed data for modeling or analysis.

% 3.2 Building the ETL Pipeline in Python

```
import pandas as pd
from sqlalchemy import create_engine
def extract_data(file_path):
   """Extract data from CSV."""
   return pd.read_csv(file_path)
def transform_data(df):
    """Data cleaning and feature engineering."""
    df.dropna(inplace=True)
    df['tenure_group'] = pd.cut(df['tenure'], bins=[0, 12, 24, 48, 72], labels=
['0-12', '12-24', '24-48', '48-72'])
    return df
def load_data(df, db_uri):
    """Load data into a SQL database."""
    engine = create_engine(db_uri)
    df.to_sql('processed_customers', engine, if_exists='replace', index=False)
# Run ETL
if name == " main ":
    raw df = extract data('data/raw/customer data.csv')
   transformed df = transform data(raw df)
    load_data(transformed_df, 'sqlite:///churn.db')
```

11 4. Case Study: End-to-End Churn Prediction Pipeline

Let's build a complete churn prediction pipeline from data ingestion to model evaluation.

4.1 Step 1: Data Ingestion

```
import pandas as pd

# Load customer data

df = pd.read_csv('data/raw/customer_data.csv')
print(df.head())
```

3 4.2 Step 2: Data Preprocessing

```
from sklearn.preprocessing import LabelEncoder

# Encode categorical variables
le = LabelEncoder()
df['gender'] = le.fit_transform(df['gender'])

# Handle missing values
df.fillna(df.mean(), inplace=True)
```

A 4.3 Step 3: Feature Engineering

```
# Feature engineering example: Monthly spending ratio
df['spending_ratio'] = df['TotalCharges'] / (df['tenure'] + 1)
```

🖲 4.4 Step 4: Model Training

```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

# Train-test split
X = df.drop(['Churn'], axis=1)
y = df['Churn']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Train model
model = RandomForestClassifier()
model.fit(X_train, y_train)

# Evaluate
predictions = model.predict(X_test)
accuracy = accuracy_score(y_test, predictions)
print(f"Model Accuracy: {accuracy:.2f}")
```

4.5 Step 5: Pipeline Automation with DVC

1. Define Stages in dvc.yaml:

```
stages:
    preprocess:
    cmd: python src/data_preprocessing.py
    deps:
        - data/raw/customer_data.csv
        - src/data_preprocessing.py
    outs:
        - data/processed/customer_data.csv

train:
    cmd: python src/model_training.py
    deps:
        - data/processed/customer_data.csv
        - src/model_training.py
    outs:
        - models/churn_model.pkl
```

2. Run the Pipeline:

```
dvc repro
```

4.6 Step 6: Visualizing Churn Insights

```
import seaborn as sns
import matplotlib.pyplot as plt

# Churn by Contract Type
sns.countplot(x='Contract', hue='Churn', data=df)
plt.title('Churn by Contract Type')
plt.show()
```

Chapter Summary

♦ Structured data projects for scalability and reusability ♦ Applied data versioning with DVC ♦ Built automated ETL pipelines for efficient workflows ♦ Developed an end-to-end churn prediction pipeline

☐ Chapter 20: Integrating Data Analysis with External Systems

Modern data analysis often extends beyond local datasets, requiring integration with **databases**, **APIs**, **cloud platforms**, and **real-time streaming systems**. This chapter guides Python experts on how to bridge the gap between traditional data analysis and scalable, production-ready data workflows.

By the end of this chapter, you will be able to:

✓ Integrate data pipelines with **SQL/NoSQL databases** and **APIs** ✓ Perform data analysis using **cloud-based data warehouses** ✓ Process **real-time data streams** with Kafka and Pandas ✓ Apply **data engineering** techniques to optimize data workflows

1. Working with Databases (SQL, NoSQL) and APIs

Seamless integration with databases and external APIs is crucial for scalable data analysis.

1.1 Connecting to SQL Databases

Libraries Used: SQLAlchemy, pandas, sqlite3, psycopg2

Example: Connecting to PostgreSQL

```
import pandas as pd
from sqlalchemy import create_engine

# PostgreSQL connection string
engine =
create_engine('postgresql+psycopg2://user:password@localhost:5432/mydatabase')

# Query data into DataFrame
df = pd.read_sql("SELECT * FROM customers", con=engine)
print(df.head())
```

Best Practices:

- Use parameterized queries to avoid SQL injection.
- Optimize queries for large datasets (use LIMIT, WHERE clauses).

1.2 Working with NoSQL Databases (MongoDB)

Libraries Used: pymongo, pandas

```
from pymongo import MongoClient
import pandas as pd

# Connect to MongoDB
client = MongoClient("mongodb://localhost:27017/")
db = client["ecommerce"]
collection = db["orders"]

# Convert MongoDB data to DataFrame
data = list(collection.find())
df = pd.DataFrame(data)
print(df.head())
```

♣ Tips:

- Use **indexes** for faster querying.
- Convert nested documents using json_normalize().

(#) 1.3 Consuming APIs for Data Extraction

Libraries Used: requests, pandas

```
import requests
import pandas as pd

# Sample API (e.g., weather data)
url = "https://api.weatherapi.com/v1/current.json"
params = {"key": "YOUR_API_KEY", "q": "London"}

response = requests.get(url, params=params)
data = response.json()

# Convert to DataFrame
df = pd.json_normalize(data)
print(df[['location.name', 'current.temp_c', 'current.condition.text']])
```

⇔ Best Practices:

- Handle rate limits and timeouts.
- Implement error handling for API failures.

2. Cloud Data Processing with AWS, BigQuery, and Snowflake

Cloud data warehouses offer scalability, security, and performance for big data processing.

2.1 Using AWS S3 for Data Storage

Libraries Used: boto3, pandas

```
import boto3
import pandas as pd
from io import StringIO

# Initialize S3 client
s3 = boto3.client('s3', aws_access_key_id='YOUR_KEY',
aws_secret_access_key='YOUR_SECRET')

# Download CSV from S3
response = s3.get_object(Bucket='mybucket', Key='data/sales.csv')
df = pd.read_csv(StringIO(response['Body'].read().decode('utf-8')))
print(df.head())
```

♣ Tips:

- Use **S3 Select** for querying large files without downloading.
- Leverage **IAM roles** for secure access.

2.2 Analyzing Data in Google BigQuery

Libraries Used: google-cloud-bigguery, pandas

```
from google.cloud import bigquery

# Initialize BigQuery client
client = bigquery.Client()

# SQL query
query = """
    SELECT customer_id, SUM(purchase_amount) as total_spent
    FROM `myproject.dataset.sales`
    GROUP BY customer_id
    ORDER BY total_spent DESC
    LIMIT 10

# Convert query result to DataFrame
df = client.query(query).to_dataframe()
print(df)
```

Best Practices:

• Optimize queries with **partitioned** and **clustered** tables.

• Use **cost estimation** before running large queries.

2.3 Snowflake for Scalable Data Warehousing

Libraries Used: snowflake-connector-python, pandas

```
import snowflake.connector
import pandas as pd
# Connect to Snowflake
conn = snowflake.connector.connect(
    user='username',
    password='password',
    account='account_name',
    warehouse='warehouse_name',
    database='database_name',
    schema='schema_name'
)
# Run query
cursor = conn.cursor()
cursor.execute("SELECT * FROM customers LIMIT 5")
df = pd.DataFrame(cursor.fetchall(), columns=[col[0] for col in
cursor.description])
print(df.head())
```

♣ Tips:

- Use **external tables** for semi-structured data (JSON, Parquet).
- Scale compute power dynamically with Snowflake Warehouses.

3. Real-Time Data Streaming with Kafka and Pandas

Real-time data analysis is essential for industries like finance, IoT, and social media analytics.

3.1 Setting Up Kafka for Streaming

Libraries Used: kafka-python, pandas

- 1. Install Kafka (local or cloud)
- 2. Start Producer and Consumer Services

3.2 Streaming Data from Kafka

```
from kafka import KafkaConsumer import json
```

```
import pandas as pd
# Initialize Kafka consumer
consumer = KafkaConsumer(
    'realtime-data',
    bootstrap_servers='localhost:9092',
    auto_offset_reset='earliest',
    value deserializer=lambda x: json.loads(x.decode('utf-8'))
)
# Stream data into DataFrame
data = []
for message in consumer:
    data.append(message.value)
    if len(data) >= 10:
        break
df = pd.DataFrame(data)
print(df.head())
```

Best Practices:

- Use consumer groups for load balancing.
- Implement backpressure handling for high-throughput streams.

4. Data Engineering for Data Science Workflows

Integrating **data engineering** practices ensures that data pipelines are efficient, maintainable, and production-ready.

1 4.1 Orchestrating Workflows with Apache Airflow

Airflow automates and manages complex workflows using Directed Acyclic Graphs (DAGs).

```
from airflow import DAG
from airflow.operators.python import PythonOperator
from datetime import datetime

def extract_data():
    print("Extracting data...")

def transform_data():
    print("Transforming data...")

def load_data():
    print("Loading data...")

with DAG('etl_pipeline', start_date=datetime(2024, 1, 1),
schedule_interval='@daily') as dag:
```

```
extract = PythonOperator(task_id='extract', python_callable=extract_data)
    transform = PythonOperator(task_id='transform',
python_callable=transform_data)
    load = PythonOperator(task_id='load', python_callable=load_data)

extract >> transform >> load
```

Best Practices:

- Use **XCom** for inter-task data sharing.
- Implement failure recovery and retries.

☑ 4.2 Building End-to-End ETL with Cloud Functions

Example Workflow:

1. Trigger: File upload to AWS S3

2. Function: AWS Lambda processes file

3. Load: Data is written to Redshift for analysis

✓ Chapter Summary

♦ Connected to **SQL/NoSQL databases** and **external APIs** ♦ Leveraged **cloud platforms** (AWS, BigQuery, Snowflake) for scalable analysis ♦ Streamed and processed **real-time data** using Kafka ♦ Implemented data engineering best practices with **Airflow** and **ETL pipelines**

oxdot Chapter 21: From Data Analysis to Data Science

As data analysis evolves into more predictive and prescriptive forms, transitioning into Data Science becomes a natural progression. This chapter bridges the gap between data analysis and machine learning by introducing essential concepts and techniques for preparing data, building models, and evaluating their performance.

By the end of this chapter, you will be able to:

- Prepare data for machine learning models <a> Use Scikit-learn for building predictive analytics pipelines
- ✓ Apply **feature engineering** with Pandas and Numpy ✓ Build and evaluate simple machine learning pipelines

1. Preparing Data for Machine Learning Models

The quality of your data directly impacts model performance. Preprocessing and feature engineering ensure that your data is suitable for machine learning algorithms.

1.1 Data Preprocessing Workflow

Key preprocessing steps:

- 1. Handling missing values
- 2. Encoding categorical variables
- 3. Scaling and normalizing features
- 4. Splitting data into training and test sets

Example: Preparing a Dataset for ML

```
import pandas as pd
from sklearn.model selection import train test split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
# Sample dataset
df = pd.DataFrame({
    'age': [25, 30, 35, None, 40],
    'income': [50000, 60000, 70000, 80000, None],
    'city': ['NY', 'LA', 'SF', 'NY', 'LA'],
    'purchased': [0, 1, 0, 1, 0]
})
# Splitting features and target
X = df.drop('purchased', axis=1)
y = df['purchased']
```

```
# Define preprocessors
numeric_features = ['age', 'income']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='mean')),
    ('scaler', StandardScaler())
])
categorical_features = ['city']
categorical_transformer = Pipeline(steps=[
    ('encoder', OneHotEncoder())
])
# Combine preprocessors
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])
# Apply preprocessing
X_preprocessed = preprocessor.fit_transform(X)
print(X_preprocessed)
```

♦ 1.2 Best Practices:

- Always scale numerical features before feeding them into distance-based models (e.g., KNN, SVM).
- Use OneHotEncoding for nominal categories and OrdinalEncoding for ordered categories.
- Impute missing values carefully—consider domain knowledge.

2. Introduction to Scikit-learn for Predictive Analytics

Scikit-learn offers a wide range of machine learning algorithms and tools for building and evaluating predictive models.

2.1 Building a Simple Classification Model

Let's build a **Logistic Regression** model to predict customer purchases.

```
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score

# Build pipeline with preprocessing and model
clf_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', LogisticRegression())
])
```

```
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Train model
clf_pipeline.fit(X_train, y_train)

# Predictions
y_pred = clf_pipeline.predict(X_test)

# Evaluate
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.2f}")
```

Material Reports Key Components:

- **Pipeline:** Combines preprocessing and modeling in one streamlined process.
- Train-Test Split: Evaluates model performance on unseen data.
- Accuracy Score: Basic evaluation metric for classification.

2.2 Evaluating Model Performance

Besides accuracy, consider these metrics:

- Precision & Recall: Useful for imbalanced datasets
- **F1 Score:** Harmonic mean of precision and recall
- ROC-AUC Curve: Measures classifier's ability to distinguish between classes

```
from sklearn.metrics import classification_report, roc_auc_score

# Detailed report
print(classification_report(y_test, y_pred))

# ROC-AUC score
y_prob = clf_pipeline.predict_proba(X_test)[:, 1]
roc_auc = roc_auc_score(y_test, y_prob)
print(f"ROC-AUC Score: {roc_auc:.2f}")
```

3. Using Pandas and Numpy for Feature Engineering

Feature engineering is the art of extracting meaningful patterns from raw data.

3.1 Creating New Features

• Date Features: Extract day, month, or weekday from timestamps

- Aggregations: Group data and calculate summaries
- Interactions: Combine existing features to capture deeper relationships

Example: Feature Engineering with Pandas

```
import pandas as pd
import numpy as np

# Sample sales data
df = pd.DataFrame({
    'date': pd.date_range(start='2024-01-01', periods=5),
    'sales': [200, 300, 250, 400, 350],
    'store': ['A', 'B', 'A', 'B', 'A']
})

# Extracting date features
df['day_of_week'] = df['date'].dt.dayofweek

# Aggregating sales by store
store_sales = df.groupby('store')['sales'].sum().reset_index()

# Creating interaction term
df['sales_per_day'] = df['sales'] / (df['day_of_week'] + 1)
print(df)
```

3.2 Handling Outliers and Skewed Data

- Log Transformations: Reduce right skew in income or price data
- **Binning:** Group continuous data into discrete buckets
- Winsorizing: Limit extreme values to reduce the effect of outliers

```
# Log transform
df['log_sales'] = np.log1p(df['sales'])

# Binning sales
df['sales_bin'] = pd.qcut(df['sales'], q=3, labels=['Low', 'Medium', 'High'])
print(df)
```

2 4. Building and Evaluating Simple Machine Learning Pipelines

Combining data preprocessing, feature engineering, and modeling into a single pipeline ensures **reproducibility** and **efficiency**.

4.1 Example: End-to-End Pipeline for Regression

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
# Sample housing dataset
housing = pd.DataFrame({
    'sqft': [1000, 1500, 2000, 2500, 3000],
    'bedrooms': [2, 3, 3, 4, 4],
    'price': [200000, 250000, 300000, 400000, 450000]
})
# Features and target
X = housing[['sqft', 'bedrooms']]
y = housing['price']
# Build pipeline
pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('model', RandomForestRegressor(n_estimators=100, random_state=42))
])
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
# Train model
pipeline.fit(X train, y train)
# Predictions and evaluation
y pred = pipeline.predict(X test)
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse:.2f}")
```

4.2 Hyperparameter Tuning with GridSearchCV

Optimize model performance using **GridSearchCV**.

```
from sklearn.model_selection import GridSearchCV

# Define parameter grid
param_grid = {
    'model__n_estimators': [50, 100, 200],
    'model__max_depth': [None, 5, 10]
}

# Apply GridSearchCV
grid_search = GridSearchCV(pipeline, param_grid, cv=3)
grid_search.fit(X_train, y_train)
```

```
# Best parameters
print("Best Parameters:", grid_search.best_params_)
```

✓ Chapter Summary

- ♦ Preprocessed and engineered features for ML models ♦ Built and evaluated pipelines using Scikit-learn
- ♦ Applied **feature engineering** with Pandas and Numpy ♦ Implemented hyperparameter tuning for model optimization

oxdot Chapter 22: Best Practices and Next Steps

As you conclude your journey through data analysis with Python, it's crucial to establish best practices that ensure clean, efficient, and scalable code. In this chapter, we'll focus on coding standards, debugging techniques, memory optimization strategies, and how to build an impressive data analysis portfolio.

By the end of this chapter, you will be able to:

✓ Write optimized and maintainable data analysis code ✓ Debug, profile, and optimize memory usage in your data pipelines <a>Maintain data quality and consistency across projects <a>Duild a professional data analysis portfolio to showcase your skills

1. Writing Clean, Efficient, and Scalable Code

Efficient code isn't just about speed; it's about readability, maintainability, and scalability. Following clean coding principles allows you to build robust data pipelines that can evolve with project needs.

1.1 Code Structuring and Modularity

- Use Functions and Classes: Encapsulate logic in reusable functions or classes.
- Follow PEP8 Guidelines: Stick to Python's standard style guide for readability.
- Leverage Type Hints: Use type annotations for clarity.

Example: Modular Data Cleaning Function

```
import pandas as pd
from typing import List
def clean dataframe(df: pd.DataFrame, columns to drop: List[str]) -> pd.DataFrame:
    Cleans the dataframe by dropping specified columns and handling missing
values.
    Args:
        df (pd.DataFrame): Input DataFrame.
        columns_to_drop (List[str]): List of columns to drop.
    Returns:
        pd.DataFrame: Cleaned DataFrame.
    df = df.drop(columns=columns_to_drop)
    df = df.dropna()
    return df
# Usage
data = {'name': ['Alice', 'Bob', None], 'age': [25, None, 30], 'city': ['NY',
'LA', 'SF']}
```

```
df = pd.DataFrame(data)
cleaned_df = clean_dataframe(df, ['city'])
print(cleaned_df)
```

4 1.2 Writing Vectorized and Pythonic Code

Avoid explicit Python loops in data-heavy tasks. Use **vectorized operations** for speed and simplicity.

```
import numpy as np

# Inefficient loop-based sum
arr = np.arange(1_000_000)
total = 0
for num in arr:
    total += num

# Vectorized sum
total_vectorized = np.sum(arr)
print(f"Vectorized Sum: {total_vectorized}")
```

Tip: Numpy and Pandas operations are optimized in C, making them significantly faster than native Python loops.

1.3 Scaling Code for Large Datasets

- Use Generators: For streaming large datasets.
- Batch Processing: Divide data into chunks using Dask or Vaex.
- Out-of-Core Computation: Use libraries like Dask for datasets that don't fit into memory.

```
import pandas as pd

# Reading large CSV in chunks
chunk_iter = pd.read_csv('large_dataset.csv', chunksize=10000)
for chunk in chunk_iter:
    print(chunk.mean())
```

🗫 2. Debugging, Profiling, and Memory Optimization

Even the cleanest code can have bottlenecks. Efficient debugging and profiling can uncover hidden inefficiencies.

2.1 Debugging Techniques

- pdb (Python Debugger): Set breakpoints and step through code.
- Logging Over Print: Use the logging library for better traceability.

```
import logging
logging.basicConfig(level=logging.INFO)
logger = logging.getLogger(__name__)

def divide(a, b):
    try:
        return a / b
    except ZeroDivisionError as e:
        logger.error("Attempted division by zero")
        return None

divide(10, 0)
```

2.2 Profiling Code Performance

Use profiling tools to pinpoint slow code sections:

- cProfile Built-in profiler for Python.
- line_profiler Profile specific functions line-by-line.
- memory_profiler Track memory usage.

```
import cProfile

def slow_function():
    total = sum([i for i in range(1_000_000)])
    return total

cProfile.run('slow_function()')
```

2.3 Memory Optimization Tips

- Use categorical data types in Pandas to reduce memory footprint.
- Apply **vectorized operations** instead of loops.
- Process large datasets using **Dask** or **Vaex**.

```
import pandas as pd
import numpy as np

# Optimize memory usage
df = pd.DataFrame({
    'category': ['A', 'B', 'C', 'A', 'B'],
```

```
'value': np.random.randint(1, 100, size=5)
})

df['category'] = df['category'].astype('category')
print(df.info())
```

III 3. Ensuring Data Quality and Consistency

Data quality is the backbone of any meaningful analysis. Inconsistent or inaccurate data can derail your insights.

3.1 Establishing Data Validation Rules

- Check for missing values and outliers.
- Validate data types and enforce schema consistency.
- Apply data integrity checks using tools like pandas_schema or Great Expectations.

```
from pandas_schema import Column, Schema
from pandas_schema.validation import CanConvertValidation, InRangeValidation

schema = Schema([
    Column('age', [CanConvertValidation(int), InRangeValidation(0, 120)]),
    Column('income', [CanConvertValidation(float)])
])

errors = schema.validate(pd.DataFrame({'age': [25, -5, 130], 'income': ['50000', '60000', 'invalid']}))
for error in errors:
    print(error)
```

♦ 3.2 Ensuring Reproducibility

- Use random seeds in all stochastic operations.
- Store all transformations and cleaning steps in scripts, not notebooks.
- Maintain a clear data lineage: track every change to the data.

4. Building Your Data Analysis Portfolio

A strong portfolio showcases your expertise and practical skills. Here's how to make yours stand out.

4.1 What to Include in Your Portfolio?

- 1. **End-to-End Projects:** Show the complete data analysis pipeline from raw data to insights.
- 2. **Diverse Domains:** Finance, healthcare, social media—show versatility.

- 3. Interactive Visualizations: Use Plotly or Dash to build engaging dashboards.
- 4. **Technical Write-Ups:** Share your thought process, challenges, and insights.

4.2 Portfolio Project Ideas

Project	Skills Demonstrated
E-commerce Sales Analysis	Data cleaning, EDA, visualization
Stock Market Trend Prediction	Time series analysis, feature engineering
Social Media Sentiment Analysis	NLP, text data wrangling, visual storytelling
Real Estate Price Prediction	Regression modeling, feature engineering
COVID-19 Data Dashboard	Real-time analytics, dashboard creation

(4.3 Sharing Your Work

- GitHub: Host code repositories and notebooks.
- **Kaggle:** Participate in competitions and share notebooks.
- **Medium/Blogs:** Write about your projects and data findings.
- LinkedIn: Share portfolio highlights and engage with the data community.

Chapter Summary

♦ Write clean, scalable, and efficient code ♦ Debug, profile, and optimize memory usage in data pipelines

♦ Implement data validation rules for consistency and accuracy ♦ Build a diverse and engaging data analysis portfolio

Appendices

The appendices provide additional resources, cheat sheets, and practical tips to solidify your understanding and improve your workflow in data analysis. Whether you're revising key concepts or looking for resources to dive deeper, this section will serve as a valuable reference.

A. Numpy, Pandas, and Matplotlib Cheat Sheets

Numpy Cheat Sheet

Operation	Example	Result	
Create array	np.array([1, 2, 3])	[1 2 3]	
Zeros/Ones array	np.zeros((2, 3))	[[0. 0. 0.], [0. 0. 0.]]	
Random numbers	np.random.rand(3, 2)	Random 3x2 array Multiplies each element by 2	
Element-wise operations	arr * 2		
Matrix multiplication	np.dot(A, B)	Matrix product of A and B	
Reshape array	arr.reshape((3, 2))	New 3x2 shaped array Selects elements 1 and 2 Elements greater than 5 Mean of array elements	
Indexing/Slicing	arr[1:3]		
Boolean indexing	arr[arr > 5]		
Aggregate functions	np.mean(arr)		

Ⅲ Pandas Cheat Sheet

Operation	Example	Result
Create DataFrame	pd.DataFrame(data)	Creates DataFrame from dict
Read CSV	<pre>pd.read_csv('data.csv')</pre>	Loads CSV into DataFrame
Select columns	df['column_name']	Series of the selected column
Filter rows	df[df['Age'] > 30]	Rows where Age > 30
Group by and aggregate	<pre>df.groupby('Dept').sum()</pre>	Aggregated data by department
Apply function	<pre>df['col'].apply(lambda x: x*2)</pre>	Applies lambda function
Merge DataFrames	pd.merge(df1, df2, on='ID')	Joins two DataFrames on 'ID'
Handle missing data	df.dropna()/df.fillna(0)	Drop or fill NaN values
Pivot tables	<pre>df.pivot_table(index='Dept')</pre>	Summarizes data by department
Save DataFrame to CSV	<pre>df.to_csv('output.csv')</pre>	Writes DataFrame to CSV file

☑ Matplotlib Cheat Sheet

```
import matplotlib.pyplot as plt
import numpy as np

# Sample data
x = np.linspace(0, 10, 100)
y = np.sin(x)

# Basic Plot
plt.plot(x, y, label='Sine Wave', color='blue')
plt.title("Sine Wave Plot")
plt.xlabel("X-axis")
plt.ylabel("Y-axis")
plt.legend()
plt.show()
```

Plot Type	Function
Line Plot	plt.plot()
Bar Chart	plt.bar()
Histogram	plt.hist()
Scatter Plot	plt.scatter()
Heatmap (Seaborn)	<pre>sns.heatmap()</pre>
Subplots	plt.subplot()
Multiple Lines	<pre>plt.plot(x, y1); plt.plot(x, y2)</pre>

B B. Performance Optimization Tips for Data Analysts

1. Numpy and Pandas Optimization

- Use Vectorized Operations: Avoid Python loops where possible.
- Minimize Memory Usage: Use astype('category') for categorical data and downcast numeric columns.
- Chunk Large Files: Read large CSV files using chunksize in pd.read_csv().

2. Code Profiling and Benchmarking

- **Use** cProfile and line_profiler for detailed performance insights.
- Leverage memory_profiler to track memory bottlenecks.

```
%load_ext memory_profiler
@profile
```

```
def memory_intensive_func():
   import numpy as np
   arr = np.random.rand(10000000)
   return np.mean(arr)
```

Ø 3. Parallel and Distributed Computing

- **Use** Dask **or** Vaex for out-of-core DataFrame manipulation.
- Apply joblib or multiprocessing for parallel computations.

O C. Common Data Analysis Pitfalls and How to Avoid Them

Pitfall	How to Avoid
Ignoring Missing Data	<pre>Use .isnull(), .dropna(), or .fillna()</pre>
Overfitting in EDA	Stick to EDA's exploratory nature before modeling
Incorrect Data Types	Check with df.dtypes and convert when needed
Misinterpreting Correlation	Correlation ≠ Causation. Use domain knowledge.
Inefficient Joins and Merges	Use indexed joins and avoid unnecessary full merges
Memory Errors with Large Files	Apply chunking, Dask, or reduce DataFrame size
Ignoring Time Zones in Time Series	Always localize timestamps using .tz_localize()

S D. Sample Datasets for Practice and Case Studies

Dataset	Source	Domain
Titanic Survival Data	Kaggle	Classification
NYC Taxi Trips	NYC Open Data	Time Series/Geo
Netflix Movies & Shows	Kaggle	EDA/Visualization
E-commerce Purchases	UCI Machine Learning Repository	Customer Behavior
Financial Market Data	Yahoo Finance API	Financial Analysis
COVID-19 Dataset	Johns Hopkins University	Time Series Analysis
Airbnb Listings	Inside Airbnb	Price Prediction

E. Resources for Further Learning (Courses, Books, and Tools)

⊟ Books

- "Python for Data Analysis" Wes McKinney (The creator of Pandas)
- "Data Science from Scratch" Joel Grus

- "Effective Pandas" Matt Harrison
- "Storytelling with Data" Cole Nussbaumer Knaflic

Online Courses

- Data Analysis with Python (Coursera) University of Michigan
- Python for Data Science (Udemy)
- Practical Data Science (DataCamp)

% Useful Tools & Libraries

- Jupyter Notebooks Interactive coding and visualization
- Seaborn Statistical data visualization
- Plotly/Dash Interactive dashboards
- Dask & Vaex Big Data handling with Pandas-like syntax
- **Great Expectations** Data validation framework

Final Thoughts

Congratulations on completing this book! So You've mastered the tools and techniques needed to perform high-level data analysis with Python.

As your next steps:

- Explore advanced domains like Machine Learning or Data Engineering
- Contribute to open-source data projects
- Build a strong portfolio and share your work with the data community

"The goal is to turn data into information, and information into insight." — Carly Fiorina

Happy Analyzing! 💋